
ER Model

CSC 343

Fall 2019

MICHAEL LIUT (MICHAEL.LIUT@UTORONTO.CA)

DEPARTMENT OF MATHEMATICAL AND COMPUTATIONAL SCIENCES

UNIVERSITY OF TORONTO MISSISSAUGA



UNIVERSITY OF
TORONTO
MISSISSAUGA



Overview of Database Design

Conceptual Designs

What are the **entities** and **relationships** in the enterprise?

What information about these entities and relationships should we store in our database?

What are the *integrity constraints* and *business rules* that hold?

i.e. We need to think about “Data Governance”

Purpose of Entity-Relationship (ER) Model



UNIVERSITY OF
TORONTO
MISSISSAUGA

Allows us to create a visual representation of the database schema design.

- These are called **entity-relationship diagrams**.
- This visual also allows us to depict some constraints imposed in our schema.

Conversion of ER designs to relational database designs.

- This will come later!



Framework for ER Model

Design is a serious business! You are the architect of the database!

Business and management know they want/need a database, but they usually don't have sufficient background or expertise to tell you what they want in it.

Sketching the key components is a great way to view the hierarchal structure, as well as an efficient way to develop a “good” working database.



Entity Sets

Entity → a “thing” or object.

Entity Set → a collection of similar entities.

- For those programmers out there; it’s similar to a class in an OO language.
- Each entity set has a **key**.

Attribute → property of an entity set.

- Attributes are simple values. e.g. int, char, or str, NOT struct, sets, etc...
- Each attribute has a **domain**.



ER Diagrams

Technicalities for drawing Diagrams (in this course):

- **Entity Set** → a rectangle.
- **Attribute** → an oval.
 - Requires a line to the oval, from the rectangle, representing its Entity Set.

This is an
attribute

NOTE: Notation varies! Some textbook represent attributes within the (entity) rectangle.




Example

Entity Set **Beers** has two attributes:

1. name
2. manf (manufacturer).

Each **Beers** entity has values for these two attributes (**entities**).
e.g. (Bud, Anheuser-Busch)





Relationships

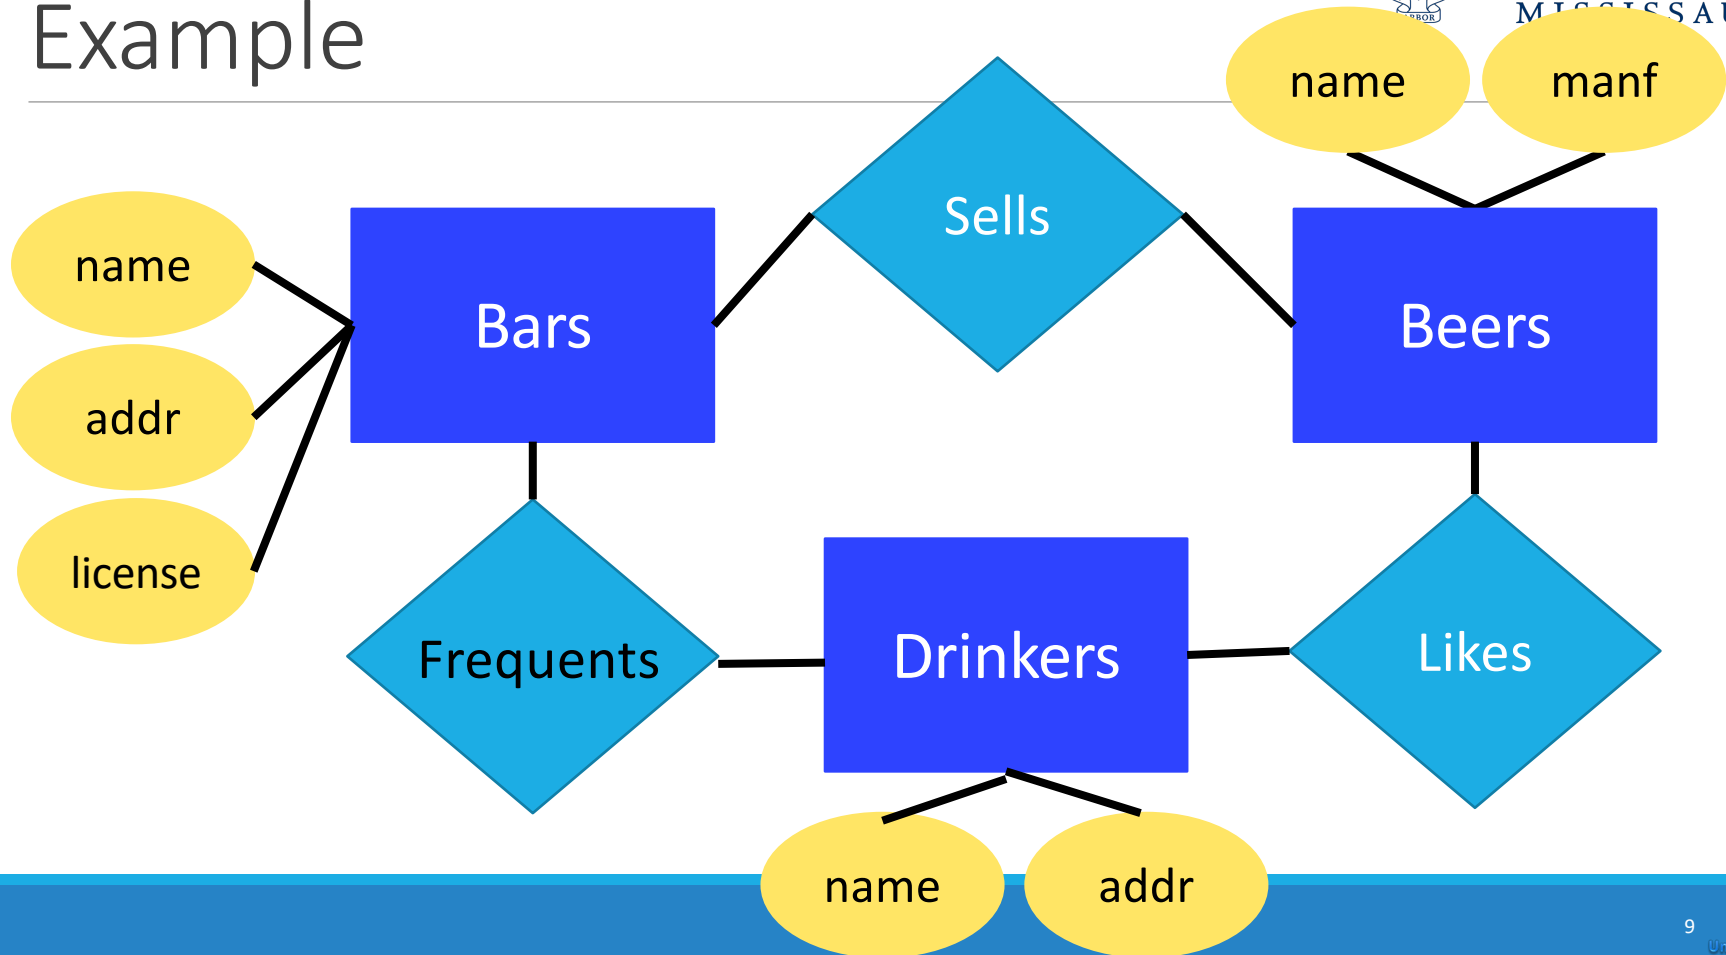
A **relationship** is a connection between two or more entity sets.

It is represented by a diamond on an E/R diagram.

- Lines connecting it to each of the entity sets involved is required.



Example





Relationship Set

The current “value” of an entity set is the set of entities that belong to it.
e.g. the set of all bars in our database.

The “value” of a relationship is a **relationship set**, a set of tuples with one component for each related entity set.

Formally:

An n -ary relationship set R relates n entity sets E_1, \dots, E_n ; where each relationship in R involves E_1, \dots, E_n .

- Same entity set could participate in different relationship sets, or different “roles” in same set.



Example

For the relationship *Sells*, we might have a relationship set like:

Bar	Beer
Joe's Bar	Canadian
Joe's Bar	Stella
Joe's Bar	Miller
Tammy's Bar	Canadian
Tammy's Bar	Corona



Multi-way Relationships

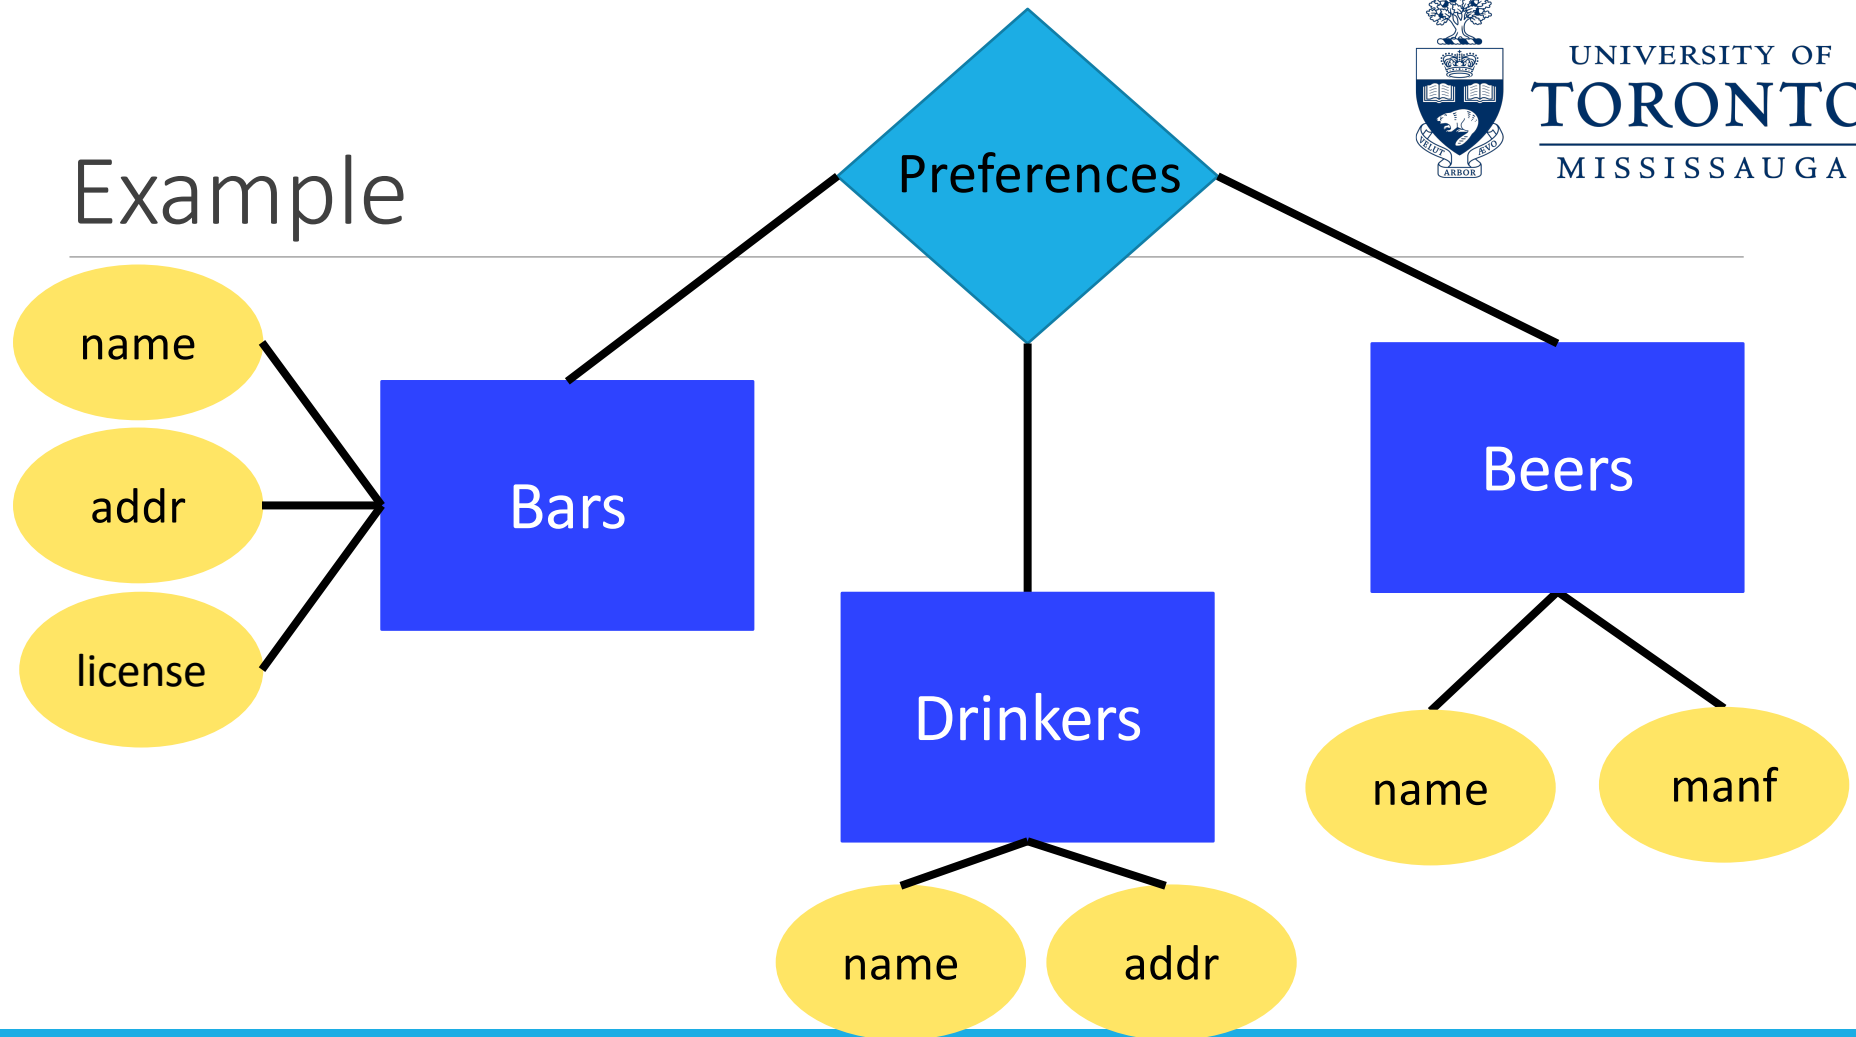
Sometimes, we need a relationship that connects more than two entity sets.

Suppose that **Drinkers** will only drink certain **Beers** at certain **Bars**.

- Our three binary relationships *Likes*, *Sells*, and *Frequents* do not allow us to make this distinction.
- But a 3-way relationship would.



Example





A Typical Relationship Set

Bar	Drinker	Beer
Joe's Bar	Jenna	Canadian
Joe's Bar	Abdi	Stella
Joe's Bar	James	Miller
Tammy's Bar	Jenna	Canadian
Tammy's Bar	Abdi	Corona
Joe's Bar	Abdi	Bud
Tammy's Bar	James	Bud
Tammy's Bar	James	Sleemans



Many-Many Relationship

Focus: binary relationships


e.g. **Sells** between **Bars** and **Beers**

In a many-many relationship, an entity of either set can be connected to many entities of the other set.

e.g. a bar sells many beers; a beer is sold by many bars.



Many-Many Illustrated



Note: each line is an instance of the binary relationship!



Many-One Relationships


Some binary relationships are many-one from one entity set to another.

Each entity of the first set is connected to at most one entity of the second set.

But an entity of the second set can be connected to zero, one, or many entities of the first set.



Many-One Illustrated



(Partial) Function on entity set.

Bars

Beers

Drinkers




Example

Favourite, from **Drinkers** to **Beers** is many-to-one

A drinker has at most one favourite beer.

A beer can be the favourite of any number of drinkers (0 included).






One-One Relationship

Each entity of either entity set is related to (at most) one entity of the other set.

e.g. Relationship **Best-Seller** between entity-sets **Manufacturers** and **Beers**.

- A beer is the best seller for $\{0|1\}$ manufacturers, and no manufacturer can have more than one best-seller (assume no ties).





Representing “Multiplicity”

A many-one relationship is depicted by an arrow entering (at most) “one” side.


A one-one relationship is depicted by an arrow entering both entity sets.

Notation:

- Rounded (open arrow) = “exactly one”
 - i.e. each entity of the first set is related to exactly one entity of the target set.



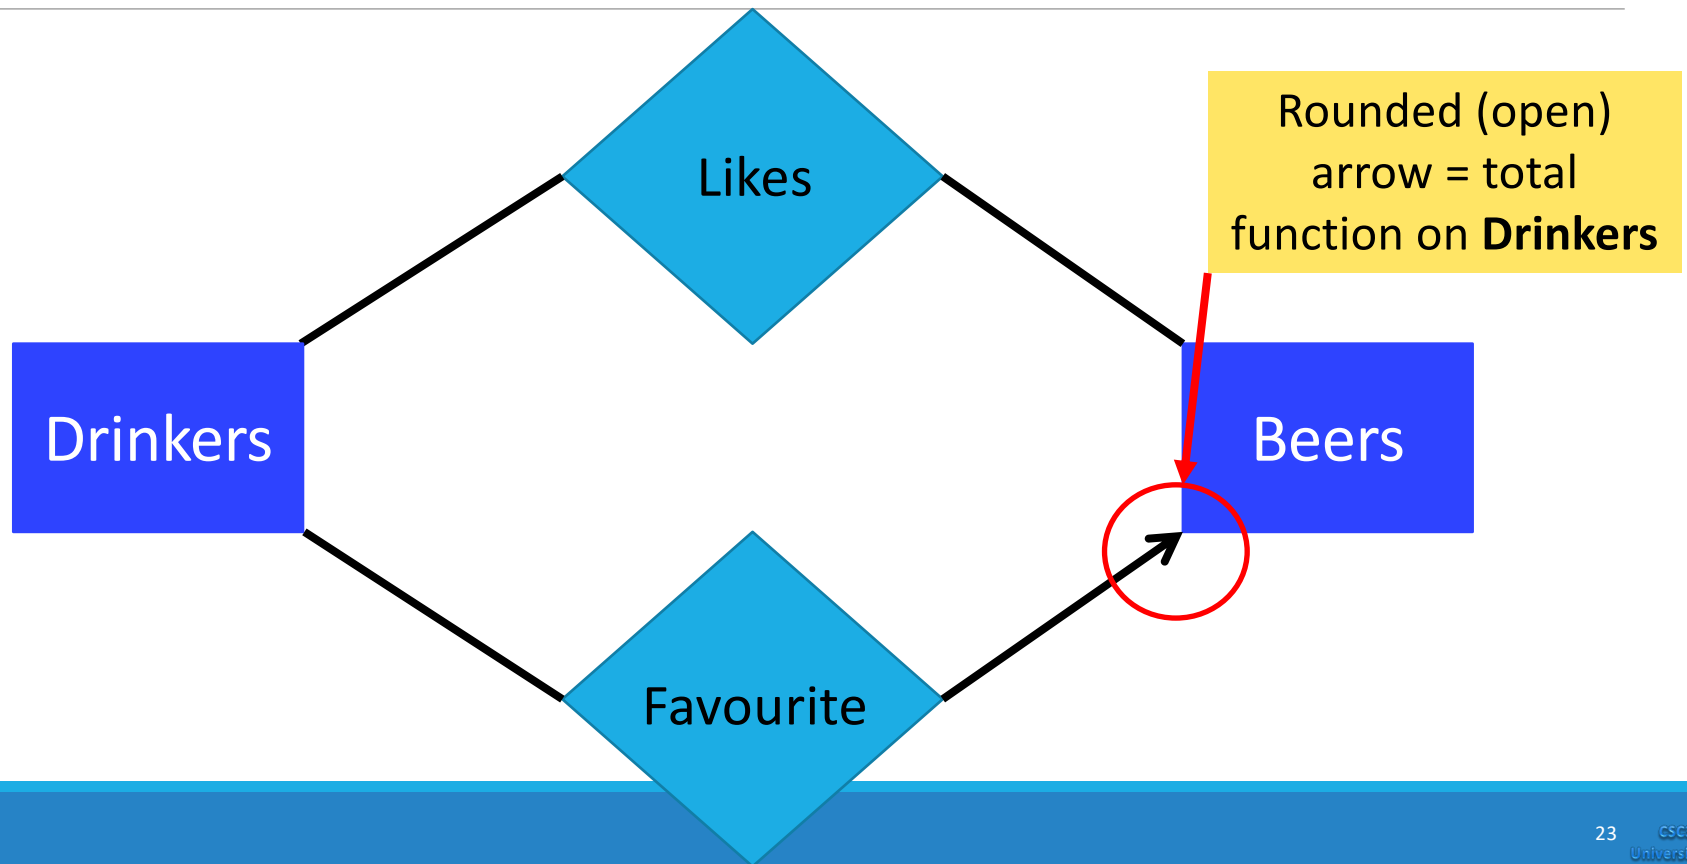
Example: Many-One Relationship



Notice that two relationships connect the same entity-sets, but are different!



Example: Many-One Relationship





Example: One-One Relationship

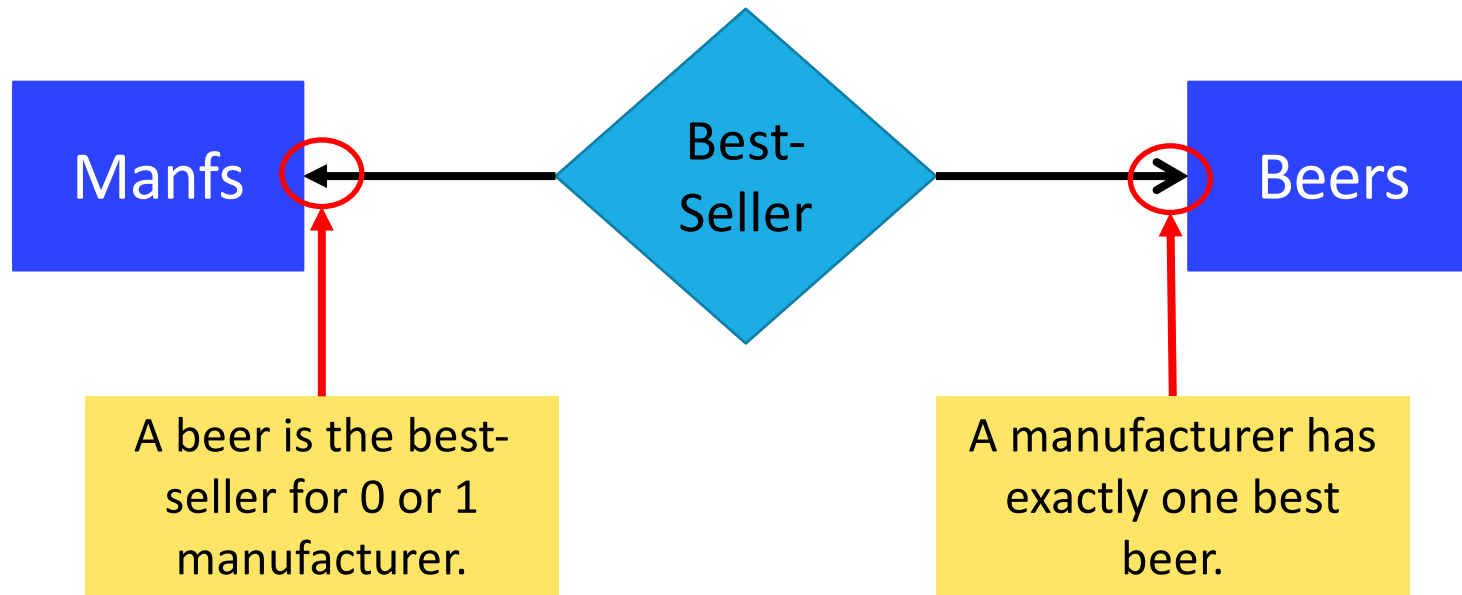
Consider **Best-Seller** between **Manfs** and **Beers**.

- Some Beers are not the best-seller of any manufacturer.
- But a Beer manufacturer has to have a best-seller.





In the E-R Diagram





Participation Constraint

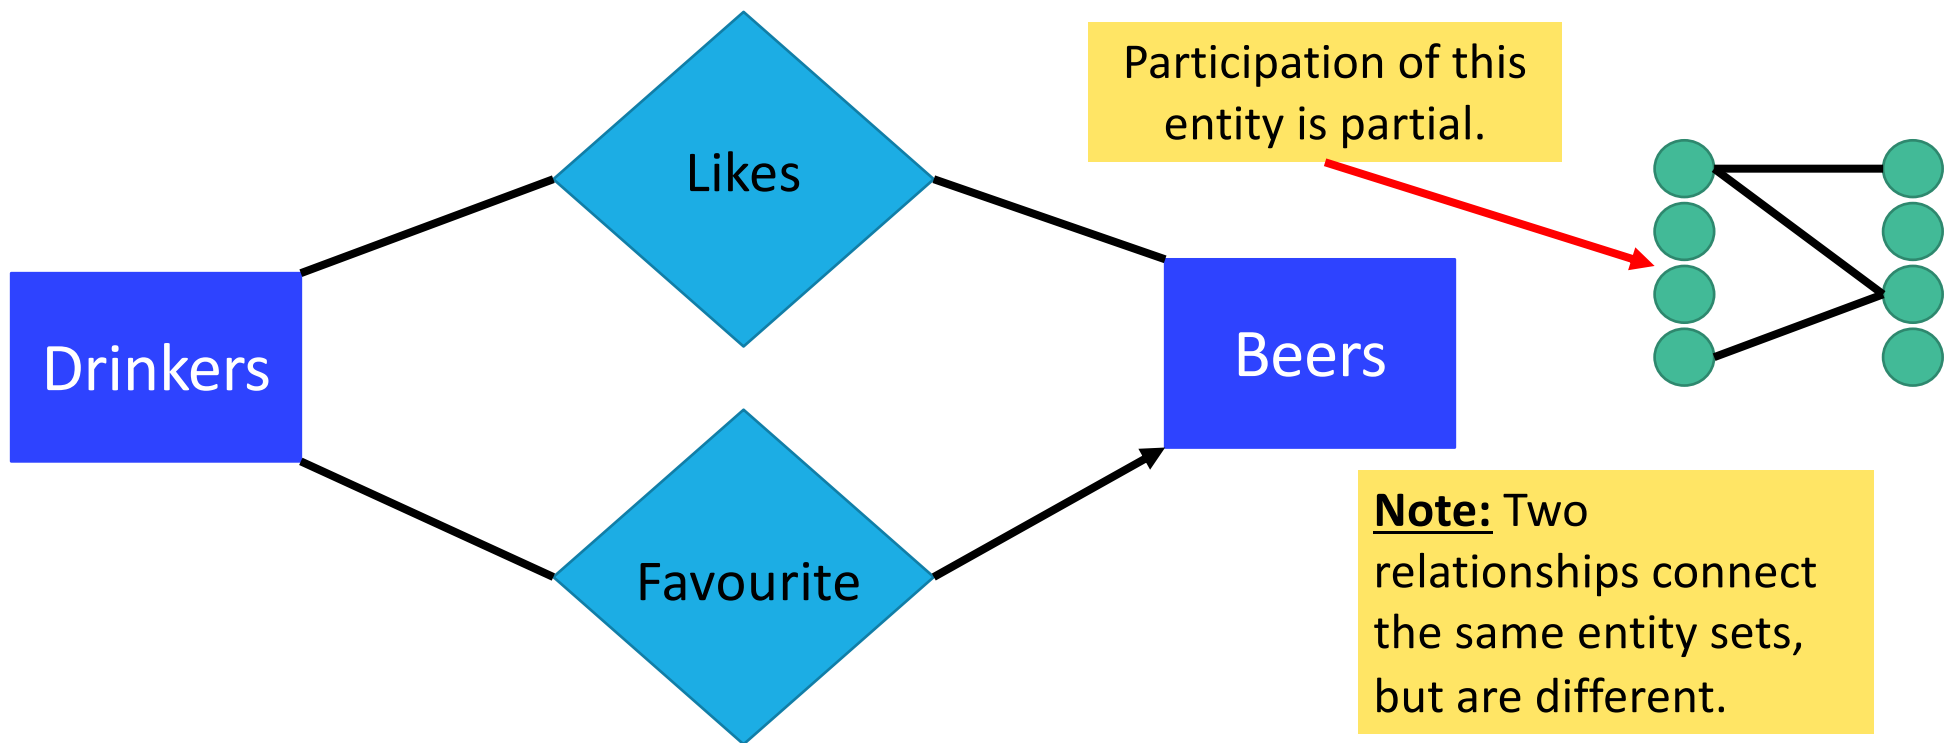
Does every student have to take a course?

- If so, this is a *participation constraint*: the participation of **Students** in **Enrolled** is said to be *total* (vs. *partial*).
- Every **SID** value in **Students** table must appear in a row of the **Enrolled** table (with a non-null **SID** value!)

Textbook Notation: total participation represented by a thick (**bolded**) line originating from entity





Example: Many-One Relationship





Example: Many-One Relationship





Participation of this
entity is total.

Rounded (open) Arrow
= total function



Alternate/Textbook Notation




Participation of this
entity is total.



Notation

Be consistent with your notation! You cannot interchange them!

Textbook



Slides





“Chen” vs. “Crow’s Foot” Notation

➤ Take a look at the attachments provided!

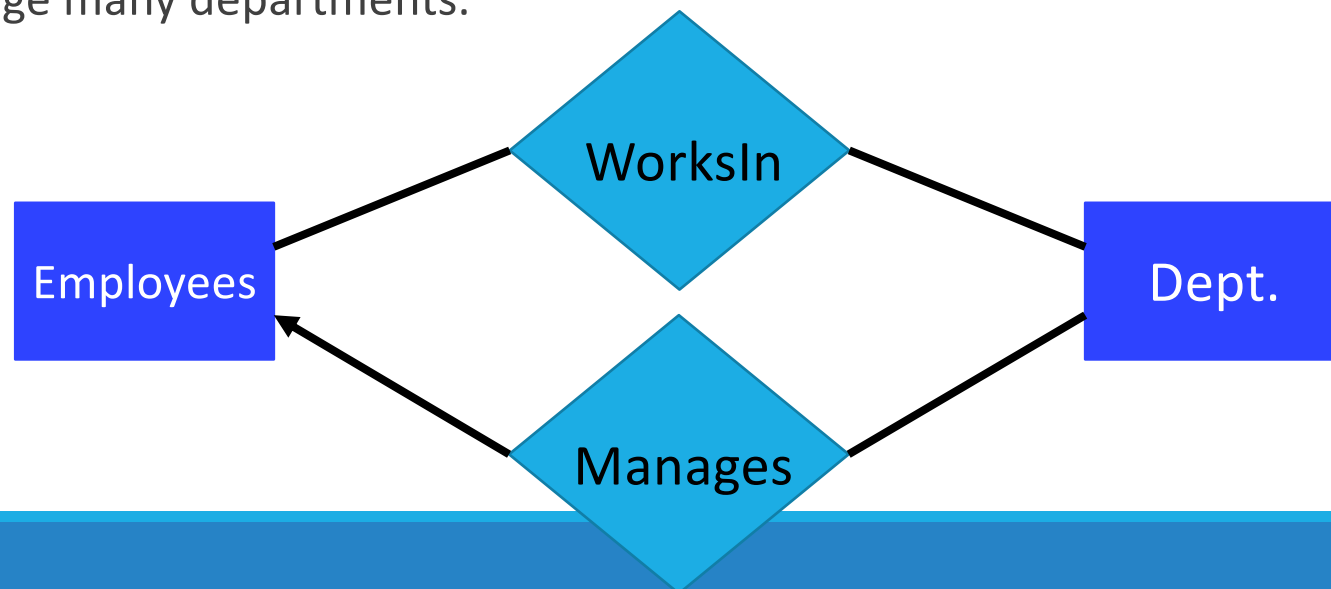
1. ERD Chen Notation
2. ERD Crow’s Notation



Key Constraints

Many-Many: “An employee can work in many departments, and a department can have many employees.”

One-Many: “A department has at most one manager, and employees can manage many departments.”






Participation Constraints

Does every department have a manager?

- If yes, then every department must appear in the manages relation:
total participation vs. **partial participation**.



—— Total participation (all)

→ Key constraint (at most one)

[textbook]



[slides]



Total participation and key constraints
(all and exactly one)



Attributes on Relationships

In certain instances, it is useful to attach an attribute to a relationship.

Think of this attribute as a property of tuples in the relationship set.


- i.e. a type of connector/bridge between entity sets to satisfy an entity.

Let's see an example!



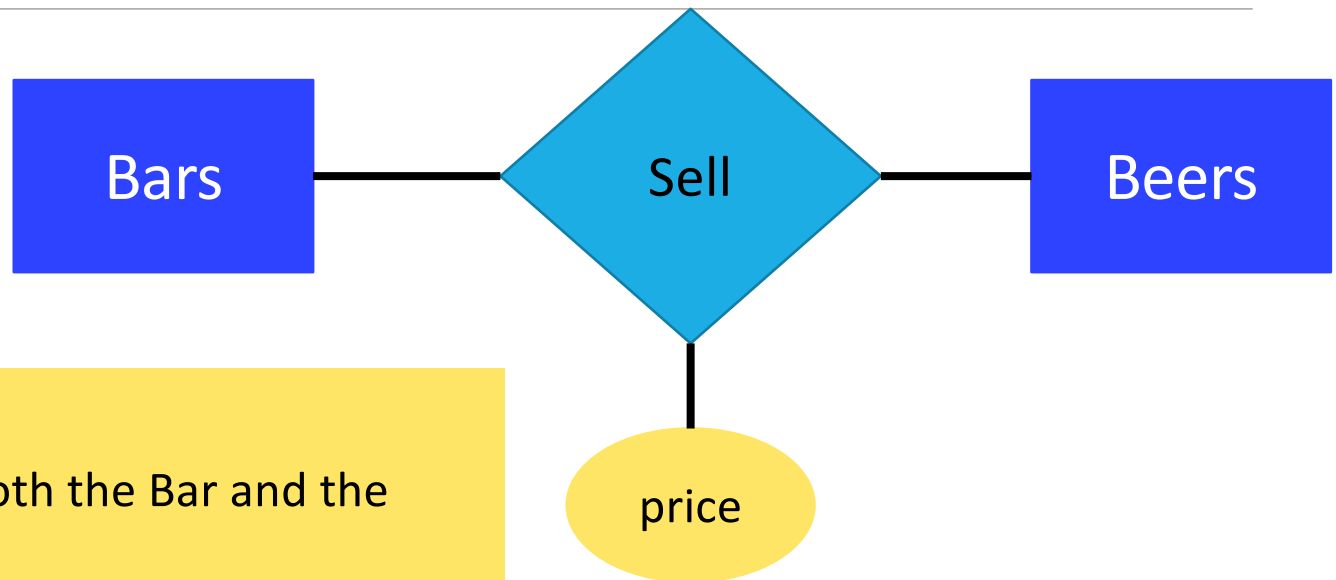
Example

Bars **sell** Beers; Beers have a **price**; Bars have a **price**.





Example



Explanation:

Price is a function of both the Bar and the Beer, not of one alone.
e.g. “The price of Miller beer at Joe’s bar.”



Equivalent Diagrams

i.e. without attributes on relationships

Create an **entity set** representing values of the **attribute (entity)**.

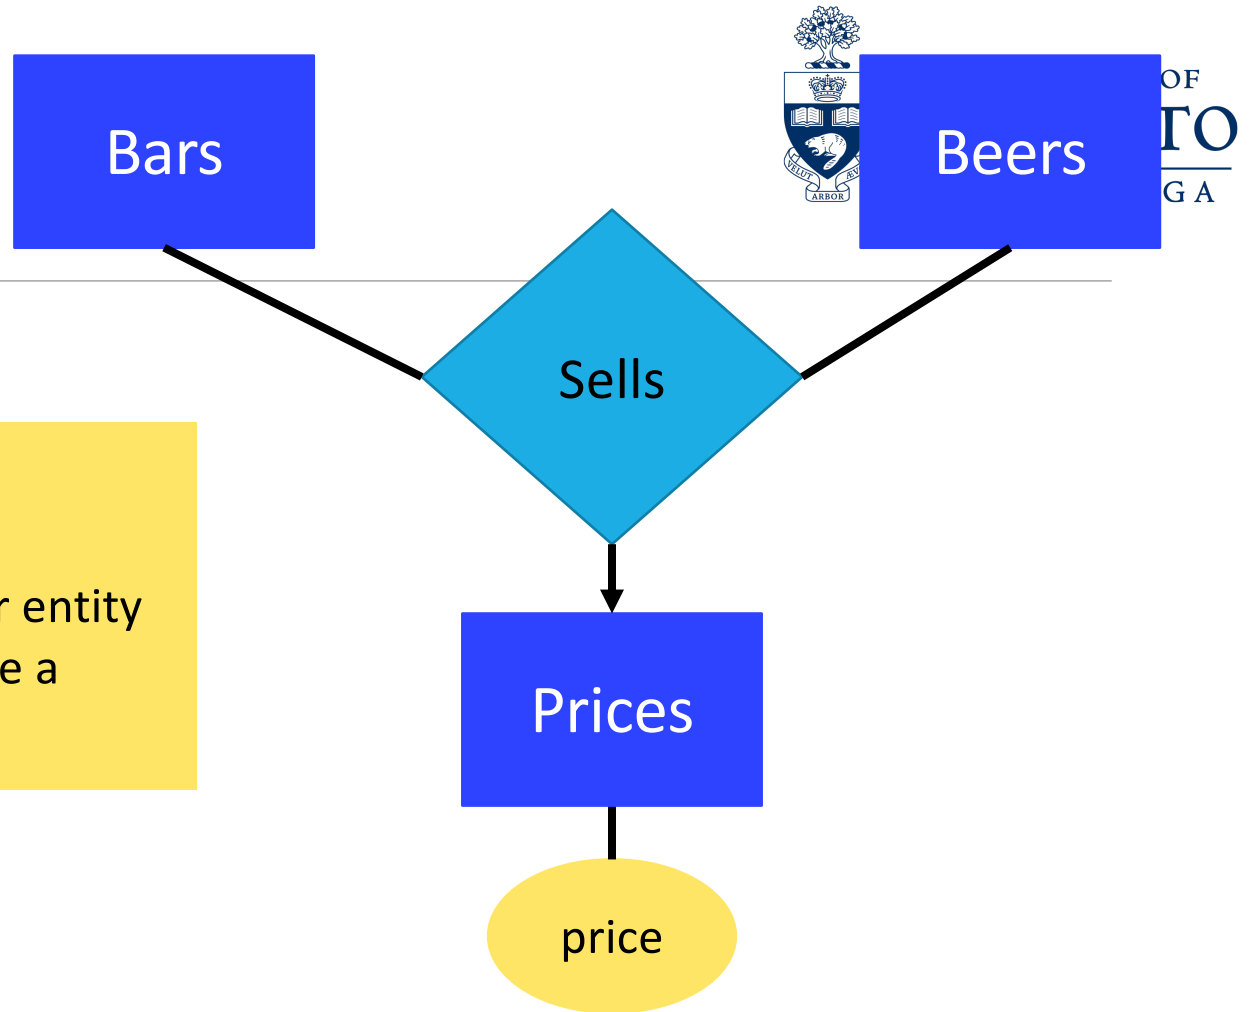
Make that **entity set** participate in the **relationship**.

Let's see an example!

Example

Notation:

Arrow from multi-way relationship = “all other entity sets together determine a unique one of these”.



Roles

Sometimes an entity set appears more than once in a relationship.

Label the edges between the relationship and entity set with names called **roles**.




EXAMPLE

Relationship Set

Husband	Wife
John	Elizabeth
Warren	Alice
...	...

husband



wife






EXAMPLE

Relationship Set

Buddy1	Buddy2
Mike	Joe
Liz	Lisa
Jenny	Peter
Courtney	Moe
...	...

1



2





Subclasses

Subclass = special case = more properties

e.g. Ales are a kind of beer.

- Not every beer is an ale, but some are.
- Let us suppose that in addition to all the *properties* (attributes and relationships) of beers also have the attribute **colour**.



Subclasses in ER Diagrams

ISA triangles indicate the subclass relationship.


- Point to the superclass.

Reasons for using **ISA**:

- To add descriptive attributes specific to a subclass.
- To identify entities that participate in a relationship.



Example



Note:

Assume subclasses form a tree.

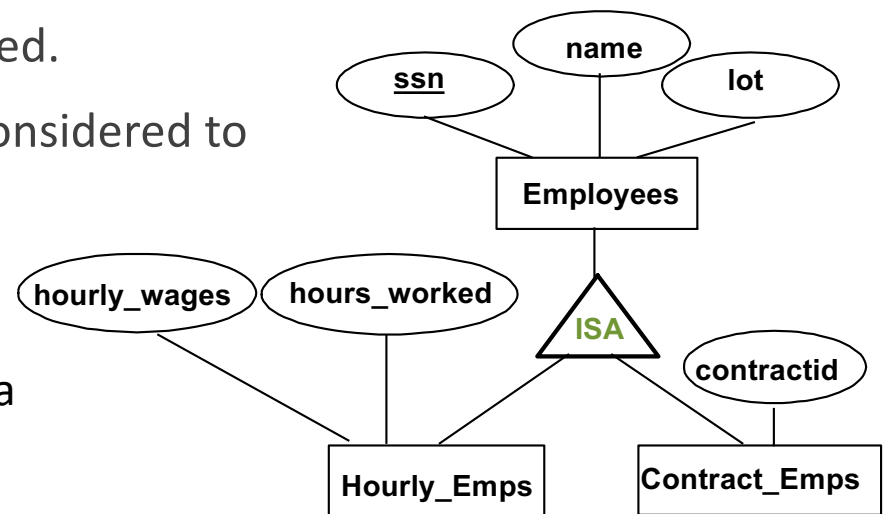


ISA (“is a”) Hierarchies

As in C++, or other PLs, attributes are inherited.

If we declare A **ISA** B, every A entity is also considered to be a B entity.

- **Overlap constraints:** Can two sub-classes contain the same entity?
e.g. Can Mike be an Hourly_Emps as well as a Contract_Emps entity?
- **Covering constraints:** Does every Employees entity have to be an Hourly_Emps or a Contract_Emps entity?

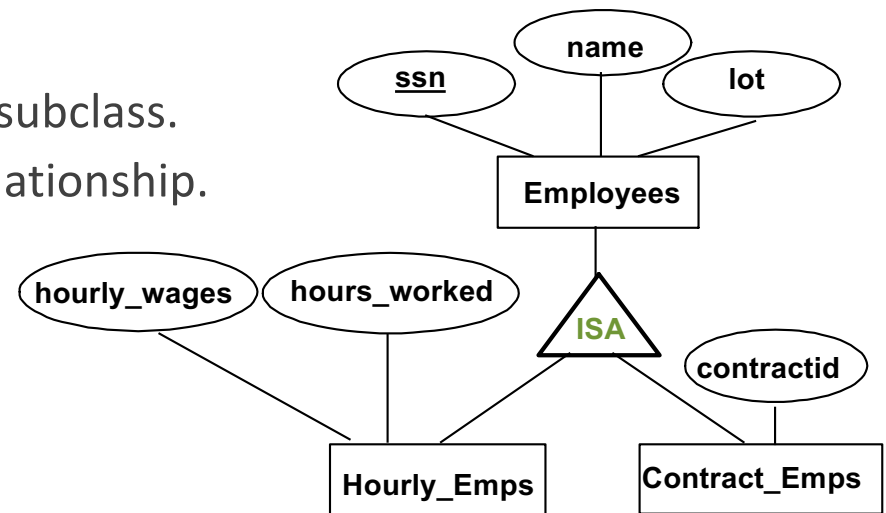




ISA (“is a”) Hierarchies

Reasons for using **ISA**:

- To add descriptive attributes specific to a subclass.
- To identify entities that participate in a relationship.

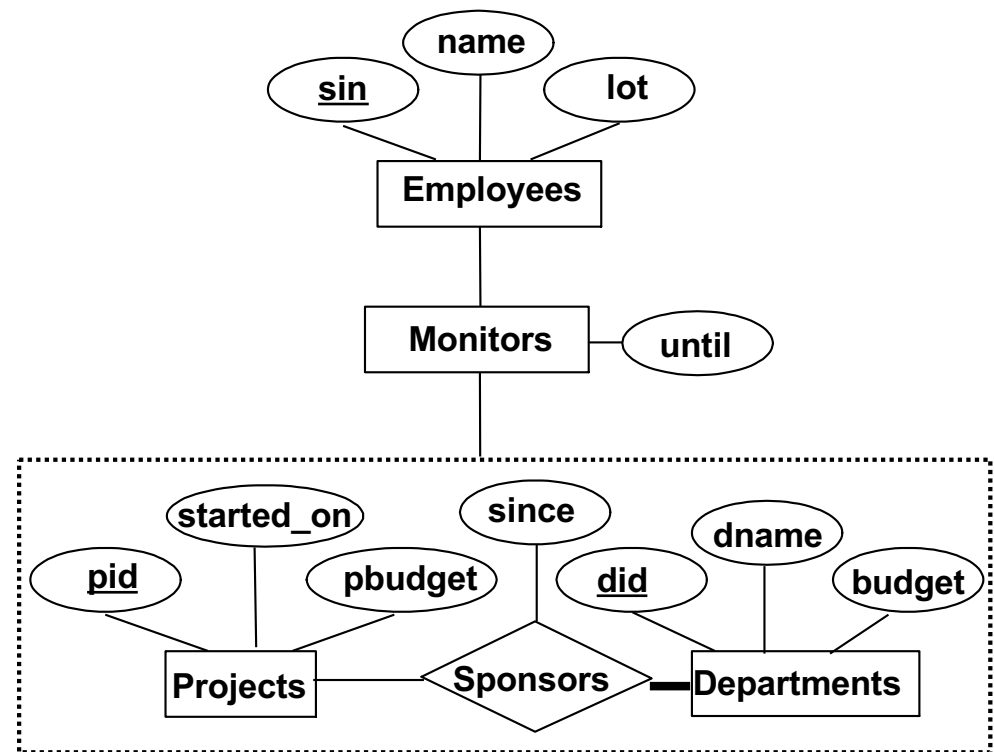




Aggregation

Used when we have to model a relationship involving (entity sets and) a relationship set.

Aggregation allows us to treat a relationship set as an entity set for the purposes of participation in (other) relationships.





Keys

A **key** is a set of attributes for one entity set such that no two entities in this set agree on all the attributes of the key.


- It is allowed for two entities to agree on some, but not all, of the key attributes.

We **must** designate a key for every entity set.

- This is identified by underlining the key attribute(s).



Example: Multi-Attribute Key




Note:

hours and **room** could also serve as a key, but we must select only one primary key (compound in this case).



Keys

In an **ISA** hierarchy, only the root entity set has a key, and it must serve as the key for all entities in the hierarchy.





Weak Entity Sets

Occasionally, entities of an entity set need “help” to identify them uniquely.

Let E represent an entity set; E is said to be **weak** if:

- E is uniquely identifiable by more than one many-one relationships from E
- and
- E includes the key of the related entities from the connected entity sets



Example: Weak Entity Set


name is almost a key for football players, but there could potentially be two players with the same name.

number is certainly not a key, since players on two teams could have the same number.

But **Teams name** and **number** (i.e. {name, number}) with relation to the player by **Plays-on** is unique.



Example: Weak Entity Set



Double diamond for supporting many-one relationship.
Double rectangle for the weak entity set.

Must be rounded because
each player needs a team
to help with the key.



Weak Entity-Set Rules

A weak entity set has one or more many-one relationships to other (supporting) entity sets.

- Not every many-one relationship from a weak entity set need be supporting.
- But supporting relationships must have a rounded arrow (entity at the “one” end is guaranteed).




Weak Entity-Set Rules

The key for a weak entity set is its own underlined attributes and the keys from supporting entity sets.

e.g. number (Players) and name (Teams) is a key for Players.



Example



e.g. **number** (**Players**) and **name** (**Teams**) is a key for **Players**.



Design Techniques

1. Avoid Redundancy.
2. Limit the use of weak entity sets.
3. Don't use an entity set when an attribute will do.



1. Avoiding Redundancy

Redundancy is saying the same thing in two (or more) different ways.


- “the inclusion of extra components that are not strictly necessary to functioning; repetition or superfluity of information.”

Redundancy wastes space and (more importantly) encourages inconsistency.

- Multiple representations of the same fact become inconsistent if we modify one and forget/do not modify its counterpart.



Example: Bad




Note:

This design states the manufacturer of a beer twice: as an attribute and as a related entity set.



Example: Good




Note:

This design gives the address of each manufacturer exactly one.



Example: Bad

Can somebody tell me why?



Note:

This design is repetitive! The manufacturer's address will be repeated for each beer. Also, in the event there are temporarily no beers for a manufacturer the value is lost.



Entity Sets vs. Attributes

An entity set should satisfy (at least) one of the following:

- It is more than the name of something (*i.e.* it has at least one non-key attribute).

OR


- It is the “many” in a many-one or many-many relationship.

Depends on the application requirements:

- If we have several addresses per employee, *address* must be an entity (since attributes cannot be set-valued).
- If the structure (city, street, etc...) is important, e.g., we want to retrieve employees in a given city, *address* must be modeled as an entity (since attribute values are atomic).



Example: Good




Manfs deserves to be an entity set because of the non-key attribute **addr**.

Beers deserves to be an entity set because it is the “many” of the many-one relationship **ManfBy**.



Example: Bad




Since the manufacturer is nothing but a name, and it is not at the “many” end of any relationship, it need not be an entity set.



Example: Good

There is no need to make the manufacturer an entity set, because we record nothing about manufacturers beside their name.





2. Limit the use of Weak Entity Sets

Novice database designers often doubt that anything could be a key by itself.

- They make all entity sets weak, supported by all other entity sets to which they are linked.

In reality, we usually create unique ID's for entity sets.

- Examples include: Social Insurance Numbers, automobile's VINs, etc...



When do we need Weak Entity Sets?

The usual reason is that there is no 'global authority' capable of creating unique IDs.

e.g. it is unlikely that there could be an agreement to assign an unique player number across all football teams in the world.



From E/R Diagrams to Relations

Entity Set \rightarrow relation

- Attributes \rightarrow attributes


Relationships \rightarrow relations whose attributes are only:

- The keys of the connected entity sets.
- Attributes of the relationship itself.



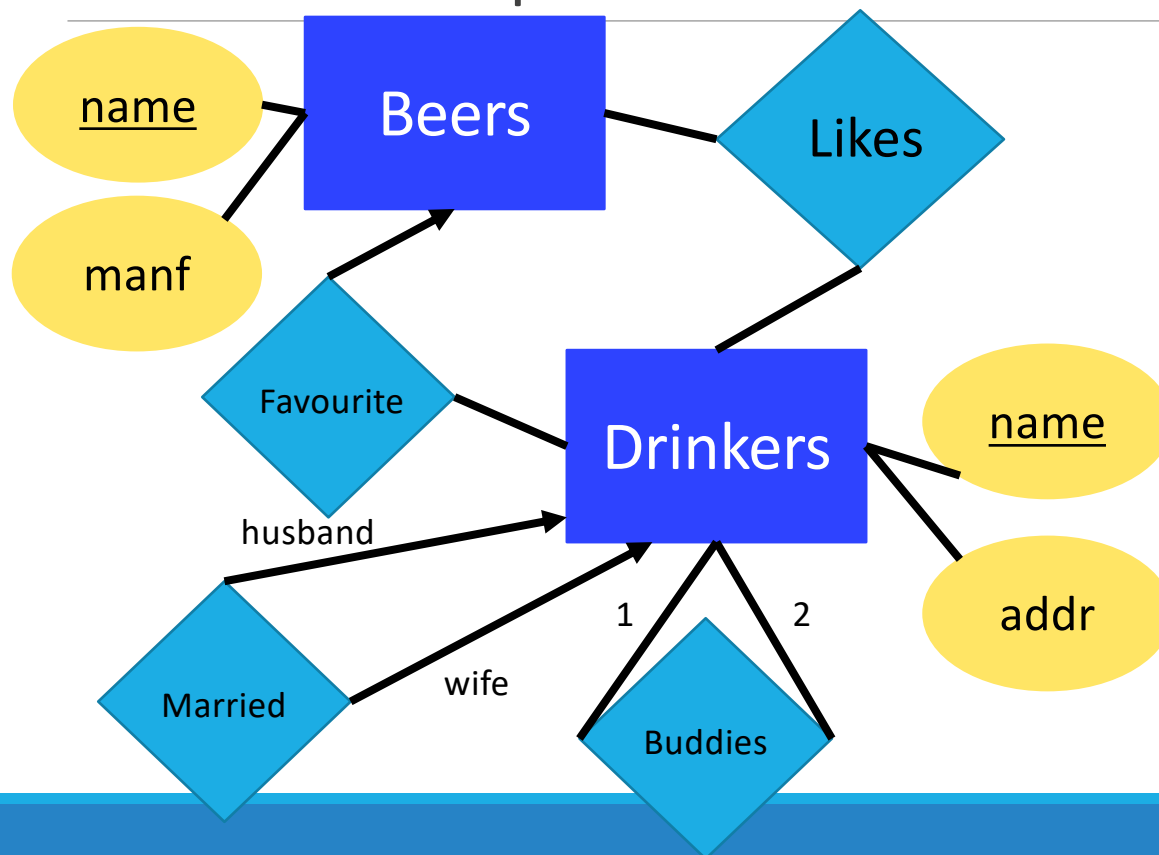
Entity Set \rightarrow Relation

Relation: Beers(name, manf)





Relationship \rightarrow Relation



Likes(drinker, beer)

Favourite(drinker, beer)

Buddies(name1, name2)

Married(husband, wife)




Handling Weak Entity Sets

Relation for a weak entity set must include attributes for its complete key (including those belonging to other entity sets), as well as its own, non-key attributes.

A supporting relationship is redundant and yields no relation (unless it has attributes).



Example: Weak Entity Set \rightarrow Relation




Hosts(hostName, location)

Logins(loginName, hostName, billTo)

At(loginName, hostName)



Example: Weak Entity Set \rightarrow Relation



Hosts(hostName, location)

Logins(loginName, hostName, billTo)

At(loginName, hostName)


Note:

At becomes part of Logins.



Entity vs. Relationship

ER diagram is OK if a manager gets a separate discretionary budget for each **Departments**.



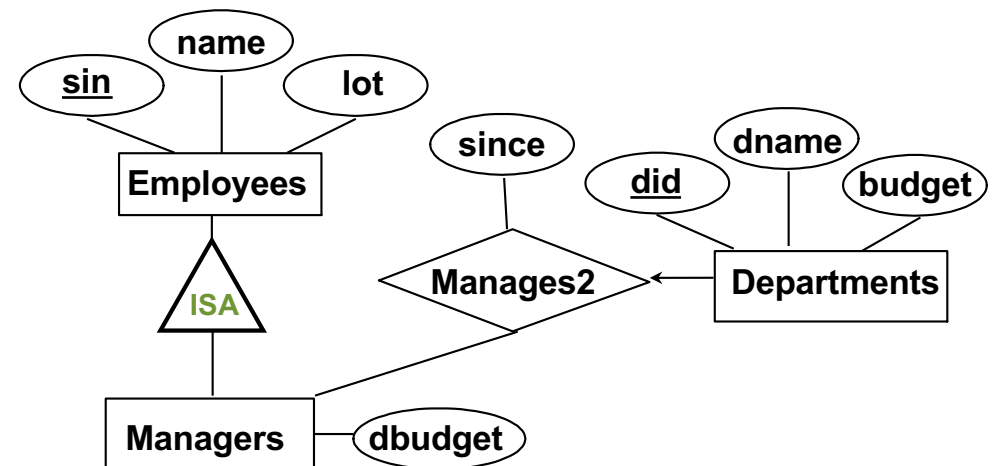
What if a manager gets a discretionary budget that covers all managed **Departments**?

Redundancy

- ❖ **dbudget** stored for each dept managed by manager.

Misleading

- ❖ Suggest **dbudget** associated with department manager combination.





Summary

Conceptual design follows requirements analysis,

- Yields a high-level description of data to be stored

ER model popular for conceptual design

- Constructs are expressive, close to the way people think about their applications.

Basic constructs: *entities*, *relationships*, and *attributes* (of entities and relationships).

Some additional constructs: *weak entities*, *ISA hierarchies*.



Summary

Several kinds of integrity constraints can be expressed in the ER model: *key constraints*, *participation constraints*, and *overlap/covering constraints* for ISA hierarchies. Some *foreign key constraints* are also implicit in the definition of a relationship set.

- Some constraints cannot be expressed in the ER model.
 - Notably, *functional dependencies* (FDs). We will talk about these later in the course.
- Constraints play an important role in determining the best database design for an enterprise.



Summary

ER design is *subjective*. There are often many ways to model a given scenario!

Analyzing alternatives can be tricky, especially for a large enterprise.

Common choices include:

- Entity vs. attribute, entity vs. relationship, binary or n-ary relationship, whether or not to use ISA hierarchies
- Ensuring good database design: resulting relational schema should be analyzed and refined further.



Flowchart Building

Online Tools:

- <https://www.draw.io>


Applications:

- OmniGraffle Pro
 - <https://www.omnigroup.com/omnigraffle>
- Microsoft PowerPoint or Word
- There are many more...

Questions?



UNIVERSITY OF
TORONTO
MISSISSAUGA



THANKS FOR LISTENING
I'LL BE ANSWERING QUESTIONS NOW



Citations, Images and Resources

Database Management Systems (3rd Ed.), Ramakrishnan & Gehrke

Some content is based off the slides of Dr. Fei Chiang - <http://www.cas.mcmaster.ca/~fchiang/>

<https://en.oxforddictionaries.com/definition/redundancy>