



---

# SQL: Data Manipulation Language

## CSC 343

### Fall 2019

---

MICHAEL LIUT ([MICHAEL.LIUT@UTORONTO.CA](mailto:MICHAEL.LIUT@UTORONTO.CA))

DEPARTMENT OF MATHEMATICAL AND COMPUTATIONAL SCIENCES  
UNIVERSITY OF TORONTO MISSISSAUGA



UNIVERSITY OF  
**TORONTO**  
MISSISSAUGA

# Why SQL?

---

SQL is a very-high-level language.

- Structured Query L- Say “what to do” rather than “how to do it.”
- Avoid a lot of data-manipulation details needed in procedural languages like C++ or Java.

DBMSs determine the “best” method of executing a query.

- This is called “query optimization”.

# Database Schemas in SQL

---

SQL is primarily a query language, used for retrieving data from a database.

- Data Manipulation Language (DML)

But SQL also includes a data-definition component for describing database schemas.

- Data Definition Language (DDL)

# Select-From-Where Statements

---

**SELECT** → desired attribute(s)

**FROM** → one or more tables (i.e. entity-sets)

**WHERE** → condition about tuples of the tables

# Recall our Example

---

Our SQL queries will be based on the following database schema.

- Underlines indicates key attributes.

Beers (name, manf)

Bars (name, addr, license)

Drinkers (name, addr, phone)

Likes (drinker, beer)

Sells (bar, beer, price)

Frequents (drinker, bar)



# Example

---

Using `Beers(name, manf)`

```
FROM Beers  
WHERE manf = 'Anheuser-Busch';
```

# Result of Query

---

name
Bud
Bud Light
Michelob
...

The answer is a relation with a single attribute, name, and tuples with the name of each beer by Anheuser-Busch, such as Bud.



UNIVERSITY OF  
**TORONTO**  
MISSISSAUGA

# Meaning of Single-Relation Query

---

Begin with the relation in the FROM clause.

Apply the selection indicated by the WHERE clause.

Apply the extended projection indicated by the SELECT clause.

# Operational Semantics - General

---

Think of a *tuple variable* visiting each tuple of the relation mentioned in FROM.

Check if the tuple assigned to the tuple variable satisfies the WHERE clause.

If so, compute the attributes or expressions of the SELECT clause using the components of this tuple.



# Operational Semantics

name	manf
Bud	Anheuser-Busch

A diagram illustrating tuple iteration over a database table. A yellow box on the left contains the text "Tuple-variable  $t$  loops over all tuples". An arrow points from this box to the first row of the table. A yellow box on the right contains the text "Check if Anheuser-Busch". An arrow points from this box to the second column of the table. Another yellow box at the top right contains the text "If so, include  $t.name$  in the result". An arrow points from this box to the first cell of the second row.



# Example

---

What beers are made by Anheuser-Busch?

```
SELECT name  
FROM Beers  
WHERE manf = 'Anheuser-  
Busch';
```

OR

```
SELECT t.name  
FROM Beers t  
WHERE t.manf = 'Anheuser-  
Busch';
```

**NOTE:** these two queries are identical.



# Asterisks ( \* ) in Select Clauses

---

When there is one relation in the **FROM** clause, \* in the **SELECT** clause stands for “all attributes of this relation”.

**Example:** Using `Beers(name, manf)`

```
SELECT *
  FROM Beers
 WHERE manf = 'Anheuser-Busch';
```

# Result of Query

---

name	manf
Bud	Anheuser-Busch
Bud Light	Anheuser-Busch
Michelob	Anheuser-Busch
...	...

The result now has each of the attributes of Beers.



# Renaming Attributes

---

If you want the result to have different attribute names, use “AS <new name>” to rename an attribute.

**Example:** Using `Beers(name, manf)`

```
SELECT name AS beer, manf  
FROM Beers  
WHERE manf = 'Anheuser-Busch';
```

# Result of Query

---

beer	manf
Bud	Anheuser-Busch
Bud Light	Anheuser-Busch
Michelob	Anheuser-Busch
...	...

The result now has each of the attributes of Beers.

# Expressions in SELECT Clauses

---

Any valid expression can appear as an element of a SELECT clause.

**EXAMPLE:** Using **Sells(bar, beer, price)**

```
SELECT bar, beer, price*80 AS priceInJPY
FROM Sells;
```



# Result of Query

---

bar	beer	priceInJPY
Joe's	Bud	285
Sue's	Miller	342
John's	Michelob	532
...	...	...

The result now has each of the attributes of Beers.



UNIVERSITY OF  
**TORONTO**  
MISSISSAUGA

# Example: Constants as Expressions

---

Using Likes(drinker, beer)

```
SELECT drinker, 'likes Bud' AS whoLikesBud
FROM Likes
WHERE beer = 'Bud';
```



# Result of Query

---

drinker	whoLikesBud
George	likes Bud
Frank	likes Bud
Jenny	likes Bud
...	...



UNIVERSITY OF  
**TORONTO**  
MISSISSAUGA

# Complex Conditions in WHERE clause

---

Boolean operators: **AND, OR, NOT**.

Comparison operators: **=, <>, <, >, <=, >=**.

# Example: Complex Condition

---

Using **Sells(bar, beer, price)**, find the price that Joe's Bar charges for Bud.

```
SELECT price  
FROM Sells  
WHERE bar = 'Joe``s Bar' AND beer = 'Bud';
```



# Patterns

---

A condition can compare a string to a pattern by:

- <Attribute> **LIKE** <pattern> or <Attribute> **NOT LIKE** <pattern>

*Pattern* is a quotes string:

- % = “any string”;
- \_\_ = “any character”.



## Example: Like

---

Using `Drinkers(name, addr, phone)`

```
SELECT name  
FROM DRINKERS  
WHERE phone LIKE '%555-_____';
```

# NULL Values

---

Tuples in SQL relations can have NULL as a value for one or more components.

The meaning is dependent on the context. In general, there are two common cases:

1. *Missing Values*: e.g., we know Joe's Bar has some address, but we don't know what it is.
2. *Inapplicable*: e.g., the value of attribute **spouse** for an unmarried person.

# Comparing NULL's to Values

---

The logic of conditions in SQL are 3 pronged:

- 1. TRUE**
- 2. FALSE**
- 3. UNKNOWN**

- Comparing any value (including NULL) with NULL yields UNKNOWN.

A tuple is in a query answer iff the **WHERE** clause is **TRUE**  
(not **FALSE** or **UNKNOWN**).



# RECALL

---

**EXAMPLE:** Using `Sells(bar, beer, price)`

```
SELECT bar, beer, price*80 AS priceInJPY  
FROM Sells;
```



# Result of Query

---

bar	beer	priceInJPY
Joe's	Bud	285
Sue's	Miller	342
John's	Michelob	532
...	...	...



UNIVERSITY OF  
**TORONTO**  
MISSISSAUGA

# String Functions in SELECT Clauses

---

**EXAMPLE:** Using `Sells(bar, beer, price)`

```
SELECT bar, INSERT(beer, 1, 0, 'EH ') AS 'Canadian beer', price
FROM Sells;
```

# Result of Query



bar	Canadian beer	price
Joe's	EH Bud	4
Sue's	EH Miller	5
John's	EH Michelob	7
...	...	...

The result reflects the renamed beers to reflect their inner Canadian.

# Three-Valued Logic

---

To understand how **AND**, **OR**, and **NOT** work in 3- valued logic

For **TRUE** results:

- **OR**: at least one operand must be **TRUE**.
- **AND**: both operands must be **TRUE**.
- **NOT**: operand must be **FALSE**.

# Three-Valued Logic

---

To understand how **AND**, **OR**, and **NOT** work in 3- valued logic

For **FALSE** results:

- **OR**: both operands must be **FALSE**.
- **AND**: at least one operand must be **FALSE**.
- **NOT**: operand must be **TRUE**.

# Three-Valued Logic

---

To understand how **AND**, **OR**, and **NOT** work in 3- valued logic

Otherwise

- The result is **UNKNOWN**.



# Example

From the following **Sells** relation:

bar	beer	price
Joe's Bar	Bud	NULL
...	...	...

```
SELECT bar  
FROM Sells  
WHERE price < 2.00 OR price >= 5.00;
```



# Multi-Relation Queries

---

Usually a combination of data from multiple (i.e. more than one) relation.

- More often than not, these are interesting queries.

We can address several relations in one query by listing them all in the **FROM** clause.

Distinguish attributes of the same name by “<relation>.<attribute>”.

# Example: Joining Two Relations

---

Using relations: Likes(drinker, beer) and Frequents(drinker, bar), find the beers likes by at least one person who frequents Joe's Bar.

```
SELECT    beer
FROM      Likes, Frequents
WHERE     bar = 'Joe''s Bar' AND Frequents.drinker = Likes.drinker;
```

# Example: Joining Two Relations

---

Alternatively, previously shown on slide #11, we can use explicit (named) tuple variables.

```
SELECT  beer
FROM    Likes l, Frequent f
WHERE   bar = 'Joe``s Bar' AND f.drinker = l.drinker;
```

# Formal Semantics

---

Almost the same for single-relation queries:

- Start with the product of all the relations in the **FROM** clause.
- Apply the selection condition from the **WHERE** clause.
- Project onto the list of attributes and expressions in the **SELECT** clause.

# Operational Semantics

---

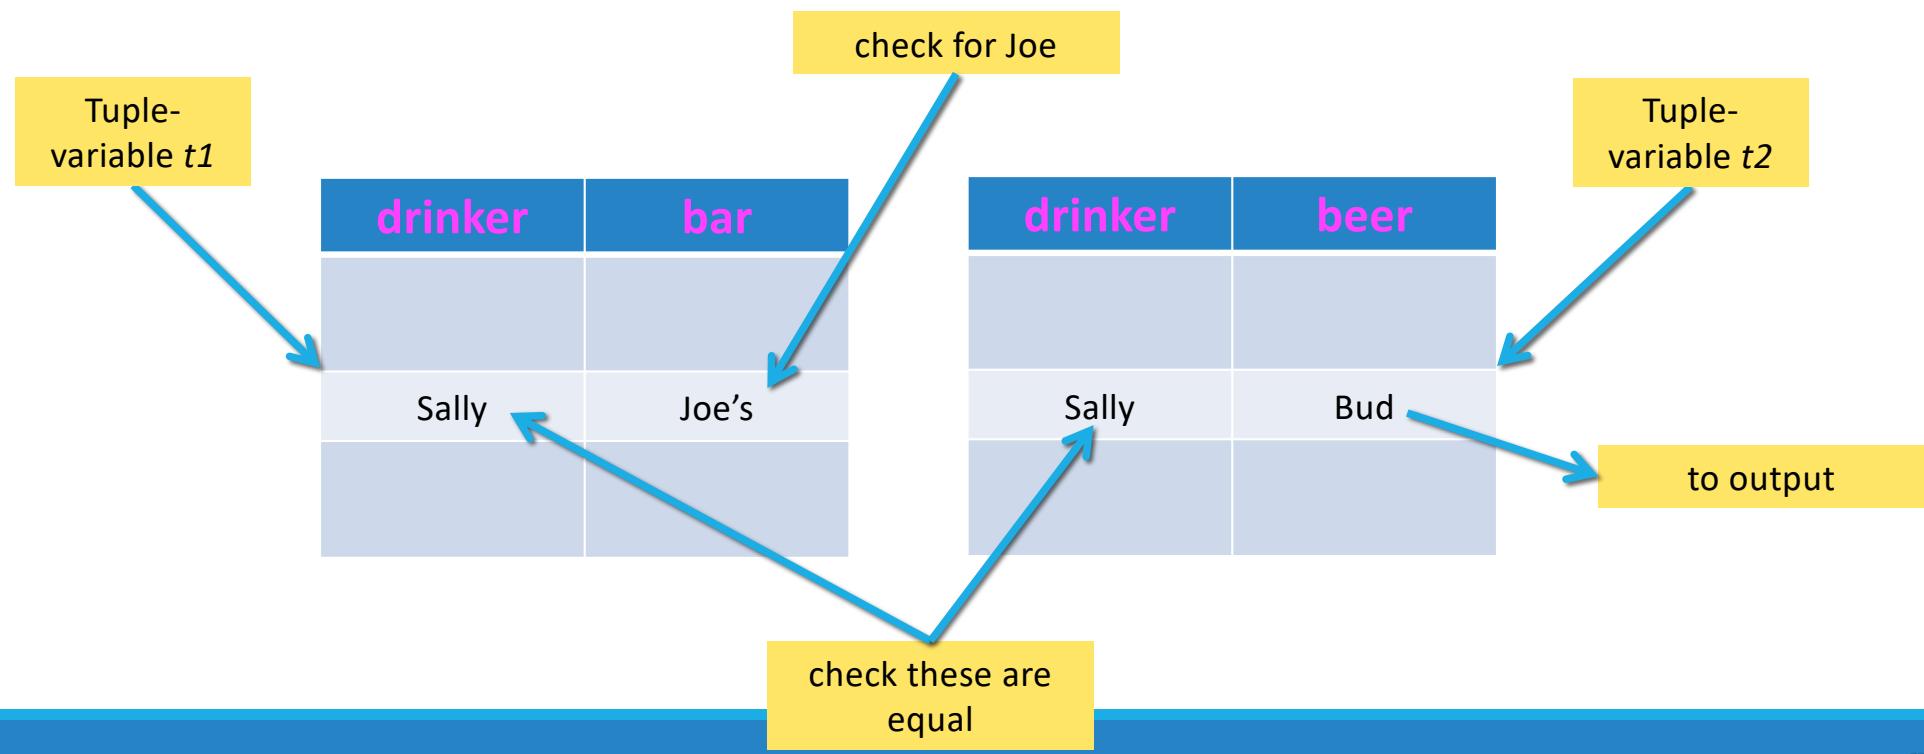
Imagine one tuple-variable for each relation in the **FROM** clause.

- These tuple-variables visit each combination of tuples, one from each relation.

If the tuple-variables are pointing to tuples that satisfy the **WHERE** clause, send these tuples to the **SELECT** clause.



# Operational Semantics



# Explicit Tuple-Variables

---

Sometimes a query needs to use two copies of the same relation.

Copies are distinguished by following the relation name by the name of a tuple-variable, in the FROM clause.

It's always an option to rename relations this way, even when not essential.

# Example: Self-Join

---

From **Beers(name, manf)**, find all pairs of beers by the same manufacturer.

- Do not produce pairs like (Bud, Bud).
- Do not produce the same pairs twice (Bud, Miller) and (Miller, Bud).

```
SELECT b1.name, b2.name
FROM Beers b1, Beers b2
WHERE b1.manf = b2.manf AND b1.name < b2.name;
```

# Sub-queries

---

A parenthesized **SELECT-FROM-WHERE** statement (subquery) can be used as a value in a number of places, including **FROM** and **WHERE** clauses.

**Example:** in place of a relation in the FROM clause, we can use a subquery and then query its result.

- To do this we must use a tuple-variable to name tuples of the result.



UNIVERSITY OF  
**TORONTO**  
MISSISSAUGA

## Example: Sub-query in FROM

Find the beers liked by at least one person who frequents Joe's Bar.

```
SELECT beer
FROM Likes, (SELECT drinker
              FROM Frequents
              WHERE bar = 'Joe``s Bar') JD
WHERE Likes.drinkers = JD.drinker;
```

Drinkers who frequent Joe's Bar





UNIVERSITY OF  
**TORONTO**  
MISSISSAUGA

# Sub-queries Often Obscure Queries

---

Find the beers liked by at least one person who frequents Joe's Bar.

```
SELECT beer
FROM Likes l, Frequents f
WHERE l.drinker = f.drinker AND bar = 'Joe``s Bar';
```

Simple Join Query

# Sub-Queries That Return One Tuple

---

If a sub-query is guaranteed to produce one tuple, then the sub-query can be used as a value.

- Usually, the tuple has one component.
- Remember SQL's 3-valued logic.

# Example: Single-Tuple Sub-query

---

Using `Sells(bar, beer, price)`, find the bars that serve Miller for the same price Joe charges for Bud.

Two queries would work:

- Find the price Joe charges for Bud.
- Find the bars that serve Miller at that price.



# Query + Sub-Query Solution

- Find the price Joe charges for Bud.
- Find the bars that serve Miller at that price.

```
SELECT bar  
FROM Sells  
WHERE beer = 'Miller' AND price =
```

**Sells(bar, beer, price)**

The price at which  
Joe sells Bud

```
(SELECT price  
FROM Sells  
WHERE bar = 'Joe ``s Bar' AND  
beer = 'Bud');
```



# Query + Sub-Query Solution

```
SELECT bar  
FROM Sells  
WHERE beer = 'Miller' AND price =
```

- Find the price Joe charges for Bud.
- Find the bars that serve Miller at that price.

Sells(bar, beer, price)

**Questions:**

1. What if the price of Bud is NULL?
2. What if the sub-query returns multiple values?

```
(SELECT price  
FROM Sells  
WHERE bar = 'Joe ``s Bar' AND  
beer = 'Bud');
```

The price at which  
Joe sells Bud

# Temporary Tables

---

Yes, they do exists! MySQL allows them!

Note: a declared temporary table must be declared!

- `CREATE TEMPORARY TABLE <temp_table_name> ...`

Temporary tables exist in:

- WHERE/FROM clauses, WITH statements, and survive for the duration a single session.
- Single session must be created in a *user temporary space*.

A few useful links for further investigation:

- <https://dev.mysql.com/doc/refman/5.7/en/create-temporary-table.html>



UNIVERSITY OF  
**TORONTO**  
MISSISSAUGA

# Recap: Conditions in WHERE Clause

---

Boolean operators: **AND, OR, NOT**.

Comparisons: `=, <>, <, >, <=, >=`.

**LIKE** operator.

SQL includes a **BETWEEN** comparison operator too.

- Let's see an example!

# Example: Between

---

Find the **names** of all **Instructors** with **salary BETWEEN \$90,000 and \$100,000.**

- i.e.  $\geq \$90,000$  and  $\leq \$100,000$ .

```
SELECT name
FROM Instructor
WHERE salary BETWEEN 90000 AND 100000;
```

# The Operator: ANY

---

$X = \text{ANY}(<\text{sub-query}>)$  is a Boolean condition that is TRUE iff  $x$  equals at least one tuple in the subquery result.

- = could be any comparison operator.

Example:  $x \geq \text{ANY}(<\text{subquery}>)$  means  $x$  is not the uniquely smallest tuple produced by the sub-query.

- NOTE: tuples must have one component only.

# The Operator: ALL

---

$x <> \text{ALL } (\text{<subquery>})$  is true iff for every tuple  $t$  in the relation,  $x$  is not equal to  $t$ .

$<>$  can be any comparison operator.

**Example:**  $x \geq \text{ALL } (\text{<subquery>})$  means there is no tuple larger than  $x$  in the sub-query result.



## Example: ALL

From Sells(bar, beer, price), find the beer(s) sold at the highest price.

```
SELECT beer
FROM Sells
WHERE price >= ALL(
    SELECT price
    FROM Sells
);
```

The price from the outer Sells  
**must not be** less than any  
price.

# The Operator: IN

---

<value> IN (<subquery>) is true iff the <value> is a member of the relation produced by the sub-query.

- Opposite: <value> NOT IN (<subquery>)

IN-expression can appear in the WHERE clauses.

WHERE <column> IN (value1, value2, ..., valueN)

# “IN” is Concise

---



```
SELECT *
FROM Cartoons
WHERE LastName
    IN ('Jetsons', 'Smurfs', 'Flinstones');
```

VS.

```
SELECT *
FROM Cartoons
WHERE LastName = 'Jetsons'
    OR LastName = 'Smurfs'
    OR LastName = 'Flinstones';
```



## Example: IN

Using `Beers(name, manf)` and `Likes(drinker, beer)`, find the name and manufacturer of each beer that Fred likes.

```
SELECT *
FROM Beers
WHERE name IN (SELECT beer
                FROM Likes
                WHERE drinker = 'Fred');
```

The set of beers that Fred likes.

# Another Example: Operator “IN” and “NOT IN”

---



Using IN and NOT IN operators with a Multiple-Row Sub-query:

- <http://www.w3resource.com/sql/subqueries/multiple-row-column-subqueries.php>



UNIVERSITY OF  
**TORONTO**  
MISSISSAUGA

## IN vs. Join

---

```
SELECT R.a  
FROM R, S  
WHERE R.b = S.b;
```

VS.

```
SELECT R.a  
FROM R  
WHERE b IN (SELECT b FROM S);
```

**Note:** IN and JOIN are different queries that can yield different results.  
Unless S.b is unique!

# IN vs. Join

---

## Equivalent to:

```
SELECT R.a
  FROM R, S
 JOIN (SELECT DISTINCT b
        FROM S
      )
  ON R.b = S.b;
```

**Performance Note:** if the joining column is not UNIQUE then IN is faster than JOIN on DISTINCT.



# IN is a Predicate About R's Tuples

SELECT R.a

FROM R

WHERE b IN (SELECT b FROM S);

One loop, over  
the tuples of R.

a	b
1	2
3	4

b	c
2	5
2	6

Two 2's.

(1,2) satisfies the  
condition; 1 is  
output once.



# This Query Pairs Tuples from R, S

```
SELECT R.a  
FROM R, S  
WHERE R.b = S.b;
```

Double loop, over  
the tuples of R  
and S.

R	
a	b
1	2
3	4

S	
b	c
2	5
2	6

(1,2) with (2,5) and (1,2)  
with (2,6) both satisfy the  
condition; 1 is output  
twice.



# Recall: Original Query

```
SELECT bar  
FROM Sells  
WHERE beer = 'Miller' AND  
      price = (SELECT price  
                FROM Sells  
                WHERE beer = 'Bud');
```

**Option 1: use IN**

**Option 2: use = ANY()**



# RECAP

---

The IN() operator is equivalent to = ANY().

For ANY(), you can use other comparison operators such as:

- >, <, etc... but this is NOT applicable for the IN() operator.
- The ability to use these operators can give you more control in your query depending on your desired output/result.

# RECAP

---

**NOTE:** the `<> ANY` operator differs from `NOT IN`:

1. `<> ANY` means `!=` (a.k.a. does not equal) any...
  - e.g. `<> ANY` means `!= a or != b or != c ...`
2. `NOT IN` means `!=` (a.k.a. does not equal) any...
  - e.g. `!= a and != b and != c ...`
3. `<> ALL` means the same as `NOT IN`.



# Example: = ANY()

Sells	bar	beer	price
	Jane	Miller	3.00
	Joe	Miller	4.00
	Joe	Bud	3.00
	Jack	Bud	4.00
	Tom	Miller	4.50

Result	bar
	Jane
	Joe

```
SELECT bar
FROM Sells
WHERE beer = 'Miller' AND price =
    ANY(SELECT price
        FROM Sells
        WHERE beer = 'Bud');
```



# The Operator: Exists

---

Exists (<subquery>) is true iff the sub-query is not empty.

**Example:** From Beers(name, manf), find the beers that are unique (i.e. only) beer made by their respective manufacturer.



## Example: Exists

```
SELECT name  
FROM Beers b1  
WHERE NOT EXISTS (
```

Set of beers with the same manf as b1, but not the same beer.

```
    SELECT *  
    FROM Beers  
    WHERE manf = b1.manf AND name <> b1.name
```

```
);
```

**Notice the scope rule:**

Manf refers to the closest nested FROM with a relation having that attribute.

**NOTE:** Some DBMSs consider this to be ambiguous.

Notice the SQL “not equals” operator.



UNIVERSITY OF  
**TORONTO**  
MISSISSAUGA

# Union, Intersection, and Difference

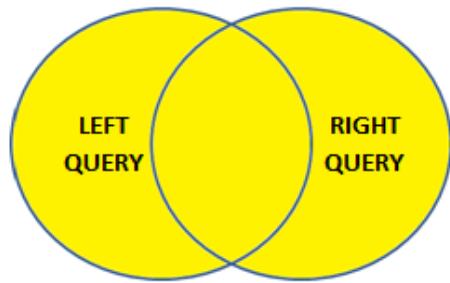
---

Union, intersection, and difference of relations are expressed by the following forms, each involving sub-queries:

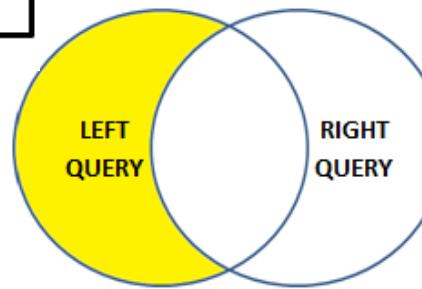
- (<subquery>) UNION (<subquery>)
- (<subquery>) INTERSECTION (<subquery>)
- (<subquery>) EXCEPT (<subquery>)



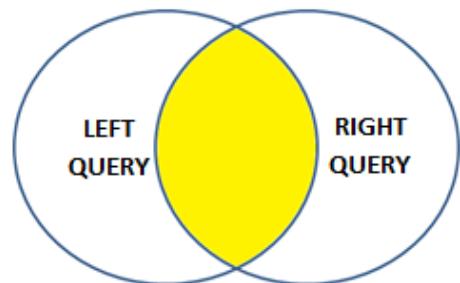
# Union, Intersection, and Difference



UNION operator returns all the unique rows from both the left and the right query. UNION ALL includes the duplicates as well



EXCEPT operator returns unique rows from the left query that aren't in the right query's results



INTERSECT operator retrieves the common unique rows from both the left and the right query

# Example: Intersection

---

Using Likes(drinker, beer) , Sells(bar, beer, price), and Frequent(drinker, bar), find the drinkers and beers such that:

1. The drinker likes the beer, and
2. The drinker frequents at least one bar that sells the beer.



# Example: Intersection – Solution

```
SELECT *  
FROM Likes
```

The sub-query is really  
a stored table!

```
INTERSECT
```

```
(SELECT drinker, beer  
FROM Sells, Frequents  
WHERE Frequents.bar = Sells.bar  
);
```

The drinker frequents a  
bar that sells the beer.

# Ordering the Display of Tuples

---

List in alphabetical order the names of all instructors

```
SELECT attribute  
FROM Table  
ORDER BY attribute <asc/desc>
```

We may specify **desc** for descending order or **asc** for ascending order, for each attribute.

**Note:** Ascending order is the default.



## Example: Order By

---

List the names of all instructors in descending alphabetical.

```
SELECT name  
FROM Instructor  
ORDER BY name;
```

vs.

```
SELECT name  
FROM Instructor  
ORDER BY name desc;
```

# Bag Semantics

---

A *bag* (or *multi-set*) is like a set, but an element may appear more than once.

Example: {1,2,1,3} is a bag.

Example: {1,2,3} is also a bag that happens to be a set.

# Bag (multi-Set Semantics)

---

SQL primarily uses bag semantics.

The SELECT-FROM-WHERE statement uses bag semantics.

- Originally for efficiency reasons.

The default for union, intersection, and difference is set semantics.

- That is, duplicates are eliminated as the operation is applied.

# Motivation: Efficiency

---

When doing projection, it is easier to avoid eliminating duplicates.

- Work on one tuple-at-a-time.

For intersection or difference, it is most efficient to sort the relations first.

- At that point you may as well eliminate the duplicates anyway.

# Control Duplicate Elimination

---

Force the result to be a set by SELECT DISTINCT...

Force the results to be a bag (i.e. don't eliminate duplicates) by the ALL operator

*i.e.* ... UNION ALL ...

## Example: DISTINCT

---

From Sells(bar, beer, price), find all the different prices charges for beers:

```
SELECT DISTINCT price  
FROM Sells;
```

Notice that without DISTINCT, each price would be listed as many times as there were bar/beer pairs at that price.

# Example: ALL

---

Using the relations: Likes(drinker, beer) and Frequents(drinker, bar):

- List drinkers who frequent more bars than they like beers, and do so as many times as the difference of those counts.

(SELECT drinker FROM Frequents)

EXCEPT ALL

(SELECT drinker FROM Likes);

# Let's Add Some HUMOUR!

---

A SQL query walks into a bar, and approaches two tables. The SQL query proceeds to ask “can I join you?”



# The Query Containment Problem: Set Semantics vs. Bag Semantics

---



First some background info...

RDBMS = Relational Database Management System.

- Created by E.F. Codd at IBM's San Jose Research Lab.
- Used to house everything from financial records to manufacturing and logistical info.
- This is based on the Relational Model (RM).

A conjunctive query is a restricted form first-order query.

- They have many desirable properties that larger classes of queries do not share.

e.g. a large class of query is relational algebra.

# The Query Containment Problem: Set Semantics *vs.* Bag Semantics

---



Relational Algebra, also created by E.F. Codd while at IBM, is a well-founded semantics for modelling the data stored in relational databases, and defining queries.

This theory is used to provide a theoretical foundation for relational databases, particularly query languages (especially SQL).

# The Query Containment Problem: Set Semantics vs. Bag Semantics



UNIVERSITY OF  
**TORONTO**  
MISSISSAUGA

Keynote talk by: Phokion Kolaitis from the University of California, Santa Cruz.

Presented in 2013 at: 7<sup>th</sup> Annual Mendelzon International Workshop on Foundation of Data Management.

**Abstract.** Query containment is a fundamental algorithmic task in database query processing and optimization. Under set semantics, the query-containment problem for conjunctive queries has long been known to be NP-complete. SQL queries, however, are typically evaluated under bag semantics and return multisets as answers, since duplicates are not eliminated unless explicitly specified. The exact complexity of the query-containment problem for conjunctive queries under bag semantics has been an outstanding and rather poorly understood open problem for twenty years. In fact, to this date, it is not even known whether conjunctive-query containment under bag semantics is decidable. The goal of this talk is to draw attention to this fascinating problem by presenting a comprehensive overview of old and not-so-old results about the complexity of the query-containment problem for conjunctive queries and their variants, under both set semantics and bag semantics.

# The Query Containment Problem: Set Semantics *vs.* Bag Semantics

---



UNIVERSITY OF  
**TORONTO**  
MISSISSAUGA

Link to conference paper: <http://ceur-ws.org/Vol-1087/keynote2slides.pdf>

Link to Hilbert's 23 problems: [https://en.wikipedia.org/wiki/Hilbert%27s\\_problems](https://en.wikipedia.org/wiki/Hilbert%27s_problems)

# Database Modifications

---

A modification command does not return a result (as a query does), but changes the database in some way.

Three types of modifications are:

1. **Insert** → inserting a tuple or tuples.
2. **Delete** → deleting a tuple or tuples.
3. **Update** → updating the value(s) of an existing tuple or tuples.

# Insertion

---

To insert a single tuple:

```
INSERT INTO <relation>
VALUES (<list_of_values>);
```

**Example:** add to Likes(drinker, beer)

```
INSERT INTO Likes
VALUES ('Sally', 'Bud');
```

# Specifying Attributes in INSERT

---

We may add to the relation name a list of attributes.

Two reasons to do so:

1. We forget the standard order of attributes for the relation.
2. We don't have values for all attributes, and we want the system to fill in missing components with NULL or a default value.



# Example: Specifying Insert Attributes

---

Another way to add the statement Sally likes Bud to Likes(drinker, beer).

```
INSERT INTO Likes (beer, drinker)  
VALUES ('Bud', 'Sally');
```



# Adding Default Values

---

In a CREATE TABLE statement, we can follow an attribute by DEFAULT and a value.

When an inserted tuple has no value for that attribute, the default value will be used.



# Example: Default Values

---

```
CREATE TABLE Drinkers (
    name CHAR(30) PRIMARY KEY,
    addr CHAR(50) DEFAULT '123 Main St.',
    phone CHAR(16)
);
```



UNIVERSITY OF  
**TORONTO**  
MISSISSAUGA

# Example: Default Values Continued

---

```
INSERT INTO Drinkers (name)  
VALUES ('Sally');
```

Result:

name	address	phone
Sally	123 Main St.	NULL



# Insert Many Tuples

---

We may insert the entire result of a query into a relation, using the form:

```
INSERT INTO <relation>
(<subquery>);
```

# Example: Insert a SubQuery

---

Using `Frequents(drinker, beer)`, enter into the new relation `Buddies(name)` all of Sally's "potential buddies".

i.e. The drinkers who frequent at least one bar that Sally also frequents.

```
INSERT INTO Buddies  
(SELECT);
```



# Result

The other drinker

INSERT INTO BUDDIES

(SELECT d2.drinker

FROM Frequent d1, Frequent d2

WHERE d1.drinker = 'Sally' AND d2.drinker <> 'Sally' AND  
d1.bar = d2.bar

);

Pairs of drinker tuples where the first is for "Sally", the second is for someone else, and the bars are the same.



# Deletion

---

To delete tuples satisfying a condition from some relation:

DELETE FROM <relation>

WHERE <condition>;



## Example: Deletion

---

Delete from Likes(drinker, beer) the tuple stating that Sally likes Bud.

```
DELETE FROM Likes  
WHERE drinker = 'Sally' AND beer = 'Bud';
```

# Example: Delete all Tuples

---

Alternatively, you can make the relation Likes empty.

```
DELETE FROM Likes;
```

**Note:** No WHERE clause is required.



# Example: Delete Some Tuples

Delete from **Beers(name, manf)** all beers for which there is another beer by the same manufacturer.

```
DELETE FROM Beers b  
WHERE EXISTS(  
    SELECT name  
    FROM Beers  
    WHERE manf = b.manf AND name <> b.name  
);
```

Beers with the same manufacturer and a different name from the name of the beer represented by tuple  $b$ .



# Semantics of Deletion

---

Suppose Anheuser-Busch makes only Bud and Bud Lite.

Suppose we come to the tuple  $b$  for Bud first.

The sub-query is non-empty, because of the Bud Lite tuple, so we delete Bud.

Now, when  $b$  is the tuple for Bud Lite, do we delete that tuple too?

**We do delete Bud Lite as well!**

Let's see how!

# Semantics of Deletion

---

Answer: Yes, we do delete Bud Lite as well.

The deletion proceeds in two stages:

1. Mark all tuples for which the WHERE condition is satisfied.
2. Delete the marked tuples.



# Updates

---

To change certain attributes in certain tuples of a relation:

UPDATE <relation>

SET <list of attribute assignments>

WHERE <condition on tuples>;



# Example: Updates

---

Change drinker Fred's phone number to 555-1234.

```
UPDATE Drinkers  
SET phone = '555-1234'  
WHERE name = 'Fred';
```

# Example: Update Several Tuples

---

Make \$4 the maximum price for beer.

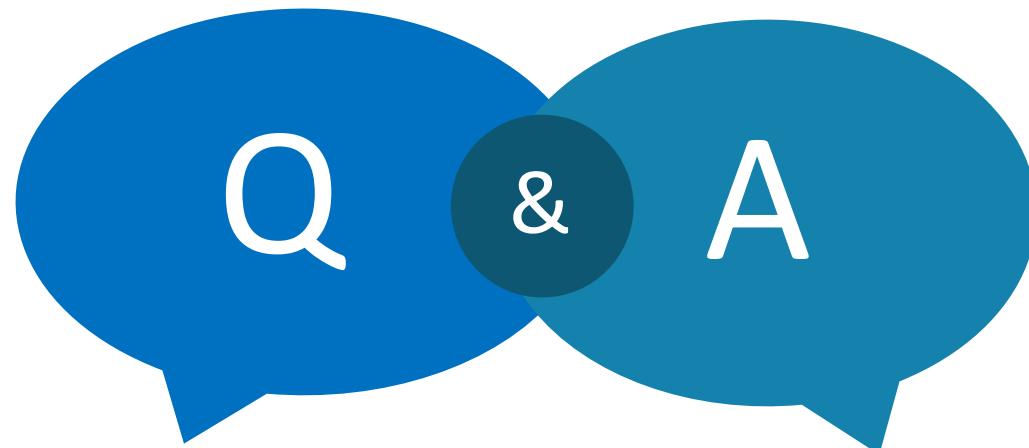
```
UPDATE Sells  
SET price = 4.00  
WHERE price > 4.00;
```

# Questions?

---



UNIVERSITY OF  
**TORONTO**  
MISSISSAUGA



THANKS FOR LISTENING  
I'LL BE ANSWERING QUESTIONS NOW



# Citations, Images and Resources

Database Management Systems (3<sup>rd</sup> Ed.), Ramakrishnan & Gehrke

Some content is based off the slides of Dr. Fei Chiang - <http://www.cas.mcmaster.ca/~fchiang/>

<http://csharpcorner.mindcrackerinc.netdna-cdn.com/UploadFile/BlogImages/06112016031910AM/sql.png>

[http://www.acoa-apeca.gc.ca/eng/investment/InvestmentHome/PublishingImages/Canada\\_flag.jpg](http://www.acoa-apeca.gc.ca/eng/investment/InvestmentHome/PublishingImages/Canada_flag.jpg)

<http://4.bp.blogspot.com/-lViQHrCxpvE/Ve3o5LdGsII/AAAAAAAAd88/Tl16JTc17Ho/s1600/difference%2Bbetween%2Bunion%2Bintersect%2Band%2Bexcept%2Bin%2Bsql%2Bserver.png>

[http://www.ibm.com/support/knowledgecenter/SSEPEK\\_10.0.0/intro/src/tpc/db2z\\_creationoftemporarytables.html](http://www.ibm.com/support/knowledgecenter/SSEPEK_10.0.0/intro/src/tpc/db2z_creationoftemporarytables.html)

[http://www.cs.newpaltz.edu/~pletcha/DB/db2\\_TempTables.html](http://www.cs.newpaltz.edu/~pletcha/DB/db2_TempTables.html)

<http://www.clipartkid.com/images/85/thumbs-up-happy-smiley-emoticon-clipart-royalty-free-public-domain-pyaRBJ-clipart.png>

<https://web.stanford.edu/~bobonich/glances%20ahead/III.logic.language.html>

[https://en.wikipedia.org/wiki/Relational\\_algebra](https://en.wikipedia.org/wiki/Relational_algebra)

[https://en.wikipedia.org/wiki/Conjunctive\\_query](https://en.wikipedia.org/wiki/Conjunctive_query)

<http://ceur-ws.org/Vol-1087/keynote2.pdf>

<http://ceur-ws.org/Vol-1087/keynote2slides.pdf>