

---

# Tutorial 10

CSC343  
Fall 2019

---

Oluwaseun Cardoso ([OLUWASEUN.CARDOSO@MAIL.UTORONTO.CA](mailto:OLUWASEUN.CARDOSO@MAIL.UTORONTO.CA))

Saihiel Bakshi ([SAIHIEL.BAKSHI@MAIL.UTORONTO.CA](mailto:SAIHIEL.BAKSHI@MAIL.UTORONTO.CA))



UNIVERSITY OF  
**TORONTO**  
MISSISSAUGA



# Recall: Definitions

---

- 1NF: No multi-valued attributes allowed.  $BCNF \subseteq 3NF \subseteq 2NF \subseteq 1NF$
- 2NF: **Non-key** attributes depend on **candidate keys**.
  - If A is a non-key attribute, then  $\exists X$  s.t.  $X \rightarrow A$ , and X is a candidate key.
- 3NF: **Non-prime** attributes depend only on **candidate keys**.
- BCNF: All non-trivial FDs have superkey LHS.



# Example: BCNF Decomposition

Drinkers(name, addr, beersLiked, manf, favBeer)

$F = \text{name} \rightarrow \text{addr}, \text{name} \rightarrow \text{favBeer}, \text{beersLiked} \rightarrow \text{manf}$

Key = name, beersLiked

- Pick BCNF violation  $\text{name} \rightarrow \text{addr}$ .
- Closure:  $\{\text{name}\}^+ = \{\text{name}, \text{addr}, \text{favBeer}\}$ .
- Decomposed relations:
  - Drinkers1(name, addr, favBeer)
  - Drinkers2(name, beersLiked, manf)



# Example: BCNF Decomposition

---

- We are not done; we need to check Drinkers1 and Drinkers2 for BCNF.
- Projecting FDs is easy here.
- For **Drinkers1(name, addr, favBeer)**, relevant FDs are **name→addr** and **name→favBeer**.
  - Thus, {**name**} is the only key and Drinkers1 is in BCNF.



# Example: BCNF Decomposition

---

- For **Drinkers2**(name, beersLiked, manf), the only FD is **beersLiked**→**manf**, and the only key is {**name**, **beersLiked**}.
  - Violation of BCNF.
- $\text{beersLiked}^+ = \{\text{beersLiked}, \text{manf}\}$ , so we decompose *Drinkers2* into:
  - **Drinkers3**(beersLiked, manf)
  - **Drinkers4**(name, beersLiked)



# Example: BCNF Decomposition

---

- The resulting decomposition of *Drinkers* :
  - *Drinkers1*(name, addr, favBeer)
  - *Drinkers3*(beersLiked, manf)
  - *Drinkers4*(name, beersLiked)
- Notice: *Drinkers1* tells us about drinkers, *Drinkers3* tells us about beers, and *Drinkers4* tells us the relationship between drinkers and the beers they like.



# Checkpoint

---

Complete the BCNF decomposition from the worksheet on quercus.

---

# Introduction to Transactions & Concurrency

---





# Transactions

---

- A sequence of many actions which are considered to be one unit of work.
  - Example:

**T1:** R(A) R(B) W(B) W(A) Commit

- R(A): Read database object A
- W(A): Writing (to) an object A
- Commit: Committing transaction
- Abort: Aborting transaction



# Schedules

- A list of actions from a set of transactions in a specific order

Example:

- **T1:** R(A) R(B) W(B) W(A) Commit
- **T2:** R(B) W(A) Commit
- **S:**  $R_1(A)$   $R_1(B)$   $R_2(B)$   $W_2(A)$   $W_1(B)$   $W_1(A)$  Commit<sub>1</sub> Commit<sub>2</sub>

S	T1	R(A)	R(B)			W(B)	W(A)	Commit	
	T2			R(B)	W(A)				Commit



# Conflict Operations

---

Two operations in a schedule are said to be conflict if they satisfy all three of the following conditions:

1. they belong to different transactions;
2. they access the same item A; and
3. at least one of the operations is a write(A).

Example in  $S_a$ :  $R1(A)$ ,  $R2(A)$ ,  $W1(A)$ ,  $W2(A)$ ,  $A1$ ,  $C2$

- $R1(A)$ ,  $W2(A)$  conflict, so do  $R2(A)$ ,  $W1(A)$ ,
- $R1(A)$ ,  $W1(A)$  do not conflict because they belong to the same transaction,
- $R1(A)$ ,  $R2(A)$  do not conflict because they are both read operations



# Write Read Conflict

1. they belong to different transactions? ✓									
2. they access the same item?									
3. at least one of the operations is a write?									
S	T1	R(A)	W(A)				R(B)	W(B)	Abort
	T2			R(A)	W(A)	Commit			
			A = X + 200	Y = A	A = Y * 1.05				
This is a <b>Write-Read</b> conflict. (dirty read)									
The problem is we are reading data that has been written by a transaction that has since aborted (and later aborted).									



# Read Write Conflict

1. they belong to different transactions? ✓
2. they access the same item? ✓
3. at least one of the operations is a write? ✓

S	T1	R(A)	R(A)		W(A)	Commit
	T2		R(A)	W(A)	Commit	
	Action	X = 5	Y = 5	A = Y - 1	A = X - 1	

This is a **Read-Write** conflict.

The problem is that after T1 has finished, T2 runs and updates A and then T1 continues to

# Write Write Conflict

1. they belong to different transactions? ✓
2. they access the same item? ✓
3. at least one of the operations is a write?

S	T1	W(A)				W(B)	Commit
	T2		W(A)	W(B)	Commit		
	Action	A = 1000	A = 2000	B = 2000		B = 1000	

This is a **Write-Write** conflict.

T2 runs during T1 running and writes to A that T1 wrote into A. The objective of T1 is to set both A and B to 1000, and the objective of T2 is to set both A and B to 2000. At the end of this schedule we have an inconsistent state.



# Serializable

- A schedule is **serializable** if the results of executing that schedule is identical to executing the transactions in the schedule in some serial order.

S	<b>T1</b>	R(A)	W(A)			R(B)	W(B)		
	<b>T2</b>			R(A)	W(A)			R(B)	W(B)

- S is serializable because it is equivalent to running, **T1; T2**;
- **T1**'s read and write of B (shaded in grey) is not affected by **T2** in S (because R(A) W(A) do not affect B).



# Any Questions?

---

- Do you have any questions?
  1. Check piazza
  2. Post the question on piazza  
(unless it's a personal question then email one of the TAs)
- If you have any content that you would like to be added in a Tutorial, please let me know by Friday!
- Email requests to:
  - [saihiel.bakshi@mail.utoronto.ca](mailto:saihiel.bakshi@mail.utoronto.ca)