
Tutorial 11

CSC343
Fall 2019

Oluwaseun Cardoso (OLUWASEUN.CARDOSO@MAIL.UTORONTO.CA)

Saihiel Bakshi (SAIHIEL.BAKSHI@MAIL.UTORONTO.CA)



UNIVERSITY OF
TORONTO
MISSISSAUGA



Transactions

- A sequence of many actions which are considered to be one unit of work.
 - Example:

T1: R(A) R(B) W(B) W(A) Commit

- R(A): Read database object A
- W(A): Writing (to) an object A
- Commit: Committing transaction
- Abort: Aborting transaction



Schedules

- A list of actions from a set of transactions in a specific order

Example:

- **T1:** R(A) R(B) W(B) W(A) Commit
- **T2:** R(B) W(A) Commit
- **S:** $R_1(A)$ $R_1(B)$ $R_2(B)$ $W_2(A)$ $W_1(B)$ $W_1(A)$ Commit₁ Commit₂

S	T1	R(A)	R(B)			W(B)	W(A)	Commit	
	T2			R(B)	W(A)				Commit



Conflict Operations

Two operations in a schedule are said to be conflict if they satisfy all three of the following conditions:

1. they belong to different transactions;
2. they access the same item A; and
3. at least one of the operations is a write(A).

Example in Sa: R1(A), R2(A), W1(A), W2(A), A1, C2

- R1(A), W2(A) conflict, so do R2(A), W1(A),
- R1(A), W1(A) do not conflict because they belong to the same transaction,
- R1(A), R2(A) do not conflict because they are both read operations



Write Read Conflict

S	T1	R(A)	W(A)				R(B)	W(B)	Abort
	T2			R(A)	W(A)	Commit			
	Action	X = A	A = X + 200	Y = A	A = Y * 1.05				

1. they belong to different transactions? ✓
2. they access the same item? ✓ (A)
3. at least one of the operations is a write? ✓

This is a **W_{rite}R_{ead}** conflict. (dirty read)

The problem is we are reading data which is not committed (and later aborted).



Read Write Conflict

S	T1	R(A)				W(A)	Commit
	T2		R(A)	W(A)	Commit		
	Action	X = 5	Y = 5	A = Y - 1		A = X - 1	

1. they belong to different transactions? ✓
2. they access the same item? ✓ (A)
3. at least one of the operations is a write? ✓

This is a **R_{ead}W_{rite}** conflict.

The problem is that after T1 has read A, T2 runs and modifies A and then T1 continues to



Write Write Conflict

S	T1	W(A)				W(B)	Commit
	T2		W(A)	W(B)	Commit		
	Action	A = 1000	A = 2000	B = 2000		B = 1000	

1. they belong to different transactions? ✓
2. they access the same item? ✓ (A)
3. at least one of the operations is a write? ✓

This is a **W_{rite}W_{rite}** conflict.

T2 runs during T1 running and overwrites the data that T1 wrote into A. The objective of T1 is to set both A and B to 1000, and the objective of T2 is to set both A and B to 2000. At the end of this schedule we have an inconsistent state.



Serializable

- A schedule is **serializable** if the results of executing that schedule is identical to executing the transactions in the schedule in some serial order.

S	T1	R(A)	W(A)			R(B)	W(B)		
	T2			R(A)	W(A)			R(B)	W(B)

- S is serializable because it is equivalent to running, **T1; T2**;
- T1**'s read and write of B (shaded in grey) is not affected by **T2** in S (because R(A) W(A) do not affect B).




Conflict Serializable

- A schedule is **conflict serializable** if it is conflict equivalent to some serial schedule.
- (two schedules are conflict equivalent if they involve the same actions of the same transactions, and the order of conflicts is the same)

S	T1	R(A)	W(A)			R(B)	W(B)		
	T2			R(A)	W(A)			R(B)	W(B)

- S is conflict serializable because the precedence/serializability graph is acyclic.






Conflict Serializable

S	T1	R(A)			W(A)	Commit		
	T2		W(A)	Commit				
	T3						W(A)	Commit

- S is **not** conflict serializable because the precedence/serializability graph is not Acyclic (cycle between T1 and T2).





View Serializable

- A schedule is **view serializable** if it is view equivalent to some serial schedule.
- For 2 schedules to be view equivalent, they must abide by the following:
 - a. If T_i reads the initial value of an object, O in S , it must also read the initial value of O , in S'
 - b. If T_i reads a value of O written by T_j in S , it must also read the value of O written by in $T_j S'$.
 - c. For each data object, the transaction, T_x that performs the final write on the object in S , must also perform the final write on the object in S' .
- If a schedule is conflict serializable, it is also view serializable. This is not true, the other way around.



View Serializable

S	T1	R(A)			W(A)	Commit		
	T2		W(A)	Commit				
	T3						W(A)	Commit

- Let's propose S' and then test it. $S' = T1; T2; T3;$
 - T1 reads the initial value of A in S. In our S' , it also reads the initial value of A.
 - No transaction reads the value of A after it is written by another transaction.
 - For the object A, T3 performs the final write on the object in S. In our S' , T3 also performs the final write.
- Therefore, S' is view serializable, even though it was not conflict serializable.



Recoverable

- In a recoverable schedule, a transaction commits only after all transactions it reads from commit.

T1	R(A)	W(A)			Commit	
T2			R(A)	W(A)		Commit

T1	R(A)	W(A)			Abort	
T2			R(A)	W(A)		Abort

Since T2 reads A from T1, it cannot commit until T1 commits. If T1 aborts, T2 has to abort as well.



Unrecoverable

T1	R(A)	W(A)				Commit
T2			R(A)	W(A)	Commit	

T2 read A from T1 but committed before T1 committed.

T1	R(A)	W(A)			Abort	
T2			R(A)	W(A)		Commit

T2 read A from T1 but T1 never committed.



Avoiding Cascading Aborts (ACA)

- In regular recoverable schedules if T2 reads A from T1, but T1 is later aborted, T2 has to be aborted as well. ACA avoids that.
- A transaction can only read data from committed transactions.

T1	R(A)	W(A)	Commit		
T2				R(A)	W(A)

- ACA \Rightarrow Recoverable (Recoverable $\not\Rightarrow$ ACA). Serial schedules are ACA and recoverable.
- This schedule is recoverable, but non ACA:

T1	R(A)	W(A)			Commit	
T2			R(A)	W(A)		Commit



Task:

Consider the two schedules below; what can you say about them? (ie: what properties to they maintain?)


- **S1:** $R_1(A) R_2(A) W_1(A) W_2(A)$
- **S2:** $R_1(A) R_2(B) W_3(A) R_2(A) R_1(B)$

Task Solution:

- S: $R_1(A)$ $R_2(A)$ $W_1(A)$ $W_2(A)$

- Serial? **No**
- Serializable? **No**
- Conflict Serializable? **No**
- View Serializable? **No**
- Recoverable? **Yes**
- ACA? **Yes**

S	T1	R(A)		W(A)	
	T2		R(A)		W(A)




Cycle!

- S: $R_1(A)$ $R_2(B)$ $W_3(A)$ $R_2(A)$ $R_1(B)$

- Serial? **No**
- Serializable? **Yes (T1;T3;T2)**
- Conflict Serializable? **Yes**
- View Serializable? **Yes**
- Recoverable? **Not sure**
- ACA? **No**

S	T1	R(A)				R(B)
	T2		R(B)		R(A)	
	T3			W(A)		



No cycle.



Any Questions?

- Do you have any questions?
 1. Check piazza
 2. Post the question on piazza
(unless it's a personal question then email one of the TAs)
- Thank you for a great semester! :)